
Differentiate Everything with a Reversible Domain-Specific Language

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Traditional machine reverse-mode automatic differentiation (AD) suffers from
2 the problem of having a space overhead that linear to time in order to trace back
3 the computational state, which is also the source of poor performance. In
4 reversible programming, a program can be executed bi-directionally, which
5 means we do not need any additional design to trace back the computational
6 state. This paper answers how practical it is to implement a programming
7 language level reverse mode AD in a reversible programming language. By
8 implementing sparse matrix operations and some machine learning applications
9 in our reversible eDSL NiLang, and benchmark the performance with
10 state-of-the-art AD frameworks, our answer is a definite positive. NiLang is an
11 open-source r-Turing complete reversible eDSL in Julia. It empowers users the
12 flexibility to tradeoff time, space, and energy rather than caching data into a
13 global tape. Manageable memory allocation makes it an excellent tool to
14 differentiate GPU kernels too.

15 1 Introduction

16 Most popular automatic differentiation (AD) packages in the market, such as TensorFlow [Abadi et al.](#)
17 [\(2015\)](#), Pytorch [Paszke et al. \(2017\)](#) and Flux [Innes et al. \(2018\)](#) implements reverse mode AD at
18 the tensor level to meet the need of machine learning. These frameworks sometimes fail to meet the
19 diverse needs in research, for example, in physics research,

- 20 1. People need to differentiate over sparse matrix operations that are important for
21 Hamiltonian engineering [Hao Xie and Wang](#), like solving dominant eigenvalues and
22 eigenvectors [Golub and Van Loan \(2012\)](#),
- 23 2. People need to backpropagate singular value decomposition (SVD) function and QR
24 decomposition in tensor network algorithms to study the phase transition problem [Golub](#)
25 [and Van Loan \(2012\)](#); [Liao et al. \(2019\)](#); [Seeger et al. \(2017\)](#); [Wan and Zhang \(2019\)](#);
26 [Hubig \(2019\)](#),
- 27 3. People need to differentiate over a quantum simulation where each quantum gate is an
28 inplace function that changes the quantum register directly [Luo et al. \(2019\)](#).

29 People have to keep adding new backward rules to the function pool. In the remaining text, we call
30 this type of AD the domain specific AD (DS-AD). To meet the diversified need, we need a general
31 purposed AD (GP-AD) too that differentiate a general program, including scalar operations
32 efficiently. Some source code transformation based AD packages like Tapenade [Hascoet and](#)
33 [Pascual \(2013\)](#) and Zygote [Innes \(2018\)](#); [Innes et al. \(2019\)](#) are close to this goal. They read the
34 source code from a user and generate a new code that computes the gradients. However, these

35 packages have their own limitations too. In many practical applications, differentiating a program
 36 might do billions of computations. Frequent caching of data slows down the program significantly,
 37 and the memory usage will become a bottleneck as well. Caching automatically for users also
 38 makes the code not compatible to GPU, it is a huge loss for the a language that supporting
 39 compiling generic codes to GPU devices like Julia [Bezanson *et al.* \(2012, 2017\)](#).

40 These needs call for a GP-AD framework that does not cache for users automatically. Hence we
 41 propose to implement the reverse mode AD on a reversible (domain-specific) programming
 42 language [Perumalla \(2013\)](#); [Frank \(2017\)](#). So that the intermediate states of a program can be
 43 traced backward with no extra effort. There have been many prototypes of reversible languages like
 44 Janus [Lutz \(1986\)](#), R (not the popular one) [Frank \(1997\)](#), Erlang [Lanese *et al.* \(2018\)](#) and
 45 object-oriented ROOPL [Haulund \(2017\)](#). These reversible languages have solid design of
 46 reversible memory management so that the memory allocation, or the time-space tradeoff is well
 47 under the programmers’ control. A reversible language has a natural trait that they can make use of
 48 reversibility so that there is no extra time or space cost to trace back a reversible operation. In
 49 machine learning, people also manage to not erasing informations that needed in the backward
 50 propagation. These neural networks includes unitary recurrent neural networks [MacKay *et al.* \(2018\)](#),
 51 normalizing flow [Dinh *et al.* \(2014\)](#), Hyperparameter learning [Maclaurin *et al.* \(2015\)](#) and
 52 residual neural networks [Behrmann *et al.* \(2018\)](#) with reversible activation functions. Utilizing
 53 reversibility is proven to decrease the memory usage by two orders in some cases, most of these
 54 applications can be written in a reversible programming language naturally without extra
 55 framework designs. Reversible programming can generalize this idea to elementary scalar
 56 operations so that programmers’ reversible thinking can help make use the reversibility more
 57 extensively to differentiate the whole programming language.

58 In the past, the primary motivation to study reversible programming is to support reversible
 59 computing devices [Frank and Knight Jr \(1999\)](#) like adiabatic complementary
 60 metal–oxide–semiconductor (CMOS) [Koller and Athas \(1992\)](#), molecular mechanical computing
 61 system [Merkle *et al.* \(2018\)](#) and superconducting system [Likharev \(1977\)](#); [Semenov *et al.* \(2003\)](#),
 62 where a reversible computing device is more energy-efficient from the perspective of information
 63 and entropy, or by the Landauer’s principle [Landauer \(1961\)](#). People tries to keep the language
 64 restrictive and well defined so that they can be compiled to future hardwares. The drawback is they
 65 can be hardly used in real computation directly, most of them do not have basic elements like
 66 floating point numbers, arrays and complex numbers that are useful in scientific computing. Not to
 67 say most of them do not have a well maintained compiler to help simulate the code on a regular
 68 device. This motivates us to build a new embeded domain specific language (eDSL) in Julia to
 69 solve these issues, so that it can be used directly to accelerate machine learning frameworks in the
 70 host language.

71 In this paper, we first introduce the language design of a reversible programming language and
 72 introduce our reversible eDSL NiLang in Sec. 2. In Sec. 3, we explain the implementation of
 73 automatic differentiation in this eDSL. In Sec. 4, we benchmark the performance of NiLang with
 74 other AD packages and explain why reversible programming AD is fast. In the appendix, we show
 75 the detailed language design of NiLang, show some examples used in the benchmark, discuss
 76 several important issues including the time-space tradeoff, reversible instructions and hardware,
 77 and finally, an outlook to some open problems to be solved.

78 2 Language design

79 2.1 A general introduction to the reversible language design

80 2.1.1 Memory management

81 A distinct feature of reversible memory management is, the content of a variable must be known
 82 when it is deallocated. We denote the allocation of a zero emptied memory as $\mathbf{x} \leftarrow \mathbf{0}$, and the
 83 corresponding deallocation as $\mathbf{x} \rightarrow \mathbf{0}$. A variable x can be allocated and deallocated in a local
 84 scope, which is called an ancilla. It can also be pushed to a stack and used later with a pop
 85 statement. This stack is similar to a traditional stack, except it zero-clears the variable after pushing
 86 and presupposes that the variable being zero-cleared before popping. Knowing the contents in the

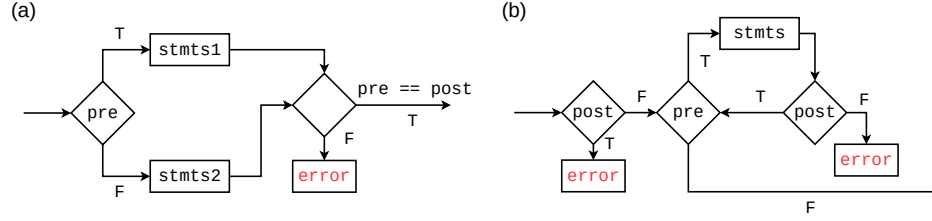


Figure 1: The flow chart for reversible (a) if statement and (b) while statement. “pre” and “post” represents precondition and postconditions respectively.

memory when deallocating is not easy. Hence Charles H. Bennett introduced the famous compute-copy-uncompute paradigm [Bennett \(1973\)](#) for reversible programming.

2.1.2 Control flows

One can define reversible if, for and while statements in a slightly different way comparing with its irreversible counterpart. The reversible if statement is shown in Fig. 1 (a). Its condition statement contains two parts, a precondition and a postcondition. The precondition decides which branch to enter in the forward execution, while the postcondition decides which branch to enter in the backward execution. After executing the specific branch, the program checks the consistency between precondition and postcondition to make sure they are consistent. The reversible while statement in Fig. 1 (b) also has two condition fields. Before executing the condition expressions, the program preassumes the postcondition is false. After each iteration, the program asserts the postcondition to be true. In the reverse pass, we exchange the precondition and postcondition. The reversible for statement is similar to irreversible ones except that after executing the loop, the program checks the values of these variables to make sure they are not changed. In the reverse pass, we exchange start and stop and inverse the sign of step.

2.1.3 Arithmetic instructions

Every arithmetic instruction has a unique inverse that can undo the changes.

- For logical operations, $y \forall= f(\text{args} \dots)$ is self reversible.
- For integer and floating point arithmetic operations, we treat $y += f(\text{args} \dots)$ and $y -= f(\text{args} \dots)$ as reversible to each other. Here f can be an arbitrary pure function such as identity, $*$, $/$ and $^$. Let’s forget the floating point rounding errors for the moment and discuss in detail in the supplementary materials.
- For logarithmic number and tropical number algebra [Speyer and Sturmfels \(2009\)](#), $y *= f(\text{args} \dots)$ and $y /= f(\text{args} \dots)$ as reversible to each other. Notice the zero element $(-\infty)$ in the Tropical algebra is not considered here.

Besides the above two types of operations, SWAP operation that exchanges the contents in two memory spaces is also widely used in reversible computing systems.

Although there are a lot reversible programming language candidates, they lack the basic components for scientific programming like arrays and complex numbers, and most of them are designed as a stand alone language that can not be embeded in other machine learning frameworks. Hence we develop an embedded domain-specific language (eDSL) NiLang in Julia language [Bezanson et al. \(2012, 2017\)](#) that implements reversible programming. One can write reversible control flows, instructions, and memory managements inside a macro. Julia is a popular language for scientific programming. We choose Julia as the host language for multiple purposes. The most important consideration is speed that crucial for a GP-AD. Its clever design of type inference and just in time compiling provides a C like speed. Also, it has a rich ecosystem for meta-programming. The package for pattern matching [MLStyle](#) allow us to define an eDSL conveniently. Last but not least, its multiple-dispatch provides the polymorphism that will be used in our AD engine. Comparing with a regular reversible programming language, NiLang features

1. array operations,

- 127 2. rich number systems, including floating point number, complex number, fixed point number
128 and logarithmic number,
- 129 3. introduce the concept of *dataview* to allow flexible data field access,
- 130 4. assuming the floating point $+=$ and $-=$ operations being reversible to each other.
- 131 5. allowing user to insert host language code like printing and asserting it as “safe”.

132 All these changes are motivated by making it a practical platform for differential applications, while
133 last two features are not compatible with reversible hardwares. By the time of writing, the version
134 of NiLang is v0.7.2. Let’s start by defining a reversible adder.

Listing 1: A reversible adder

```
135 @i function adder(y!::Real, x::Real)
    y! += x
end

@assert adder(2, 3) == (5, 3)
@assert (~adder)(5, 3) == (2, 3)
```

136 Macro `@i` generates two functions that reversible to each other `adder` and `~adder`, each defines a
137 mapping $\mathbb{R}^2 \rightarrow \mathbb{R}^2$. The `!` after a symbol is a part of the name to indicate that a variable is changed. A
138 reversible `+=` instruction is always defined as `y += f(args...)`, where `f` is a mapping that allows
139 to be irreversible, or just leave empty for identity mapping. We can easily check these two functions
140 are reversible to each other. Then let’s see a more advanced example of computing the complex
141 valued log (a built in function).

Listing 2: Reversible complex valued log function `y += log(|x|) + iArg(x)`.

```
142 @i @inline function (:+=)(log)(y!::Complex{T}, x::Complex{T}) where T
    @routine begin
        n ← zero(T)
        n += abs(x)
    end
    y!.re += log(n)
    y!.im += angle(x)
    ~@routine
end
```

143 Here, the macro `@inline` tells the compiler that this function can be inlined. `n ← zero(T)` is the
144 ancilla allocation statement. One can input “ \leftarrow ” and “ \rightarrow ” by typing “`\leftarrow`[TAB KEY]” and
145 “`\rightarrow`[TAB KEY]” respectively in a Julia editor or REPL. `@routine` and `~@routine` are
146 macros for computing and uncomputing. i.e. `~@routine` means running the statement marked with
147 `@routine` backwards. One can use the `begin ... end` statement to wrap multiple statements as
148 one. NiLang view every field of a variable as mutable, so that the real part (`y!.re`) and imaginary
149 (`y!.im`) of a complex number can also be changed directly.

150 We can want to apply this log function to an array, we can define

Listing 3: Applying the log function to an array.

```
151 @i function broadcasted_log!(y!::Array{Complex{T}, N}, x::Array{Complex{T}, N}) where T
    N ← min(length(x), length(y!))
    for i=1:N
        y![i] += log(x[i])
    end
end
```

152 3 Reversible automatic differentiation

153 3.1 First order gradient

154 The instructions executed by the reversible program looks like

Listing 4: The expanded function body of Listing. 3.

```
155 N ← min(length(x), length(y!))
    for i=1:N
      @routine begin
        nsq ← zero(T)
        n ← zero(T)
        nsq += x[i].re ^ 2
        nsq += x[i].im ^ 2
        n += sqrt(nsq)
      end
      y![i].re -= log(n)
      y![i].im += atan(x[i].im, x[i].re)
    ~@routine
  end
  N → min(length(x), length(y!))
```

Listing 5: The inverse of Listing. 4.

```
N ← min(length(x), length(y!))
for i=N:-1:1
  @routine begin
    nsq ← zero(T)
    n ← zero(T)
    nsq += x[i].re ^ 2
    nsq += x[i].im ^ 2
    n += sqrt(nsq)
  end
  y![i].re -= log(n)
  y![i].im -= atan(x[i].im, x[i].re)
~@routine
end
N → min(length(x), length(y!))
```

156 Then we insert

Listing 6: Insert the gradient code into Listing. 5.

```
157 N ← min(length(x), length(y!))
    for i=N:-1:1
      @routine begin
        nsq ← zero(GVar{T,T})
        n ← zero(GVar{T,T})

        gsqa ← zero(T)
        gsqa += x[i].re.x * 2
        x[i].re.g -= gsqa * nsq.g
        gsqa -= nsq.x * 2
        gsqa -= x[i].re.x * 2
        gsqa → zero(T)
        nsq.x += x[i].re.x ^ 2

        gsqb ← zero(T)
        gsqb += x[i].im.x * 2
        x[i].im.g -= gsqb * nsq.g
        gsqb -= x[i].im.x * 2
        gsqb → zero(T)
        nsq.x += x[i].im.x ^ 2

        @zeros T ra rb
        ra += sqrt(nsq.x)
        rb += 2 * ra
        nsq.g -= n.g / rb
        rb -= 2 * ra

        ra -= sqrt(nsq.x)
        ~@zeros T ra rb
        n.x += sqrt(nsq.x)
      end

      y![i].re.x -= log(n.x)
      n.g += y![i].re.g / n.x

      y![i].im.x -= atan(x[i].im.x, x[i].re.x)
      @zeros T xy2 jac_x jac_y
      xy2 += abs2(x[i].re.x)
      xy2 += abs2(x[i].im.x)
      jac_y += x[i].re.x / xy2
      jac_x += (-x[i].im.x) / xy2
      x[i].im.g += y![i].im.g * jac_y
      x[i].re.g += y![i].im.g * jac_x
      jac_x -= (-x[i].im.x) / xy2
      jac_y -= x[i].re.x / xy2
      xy2 -= abs2(x[i].im.x)
      xy2 -= abs2(x[i].re.x)
      ~@zeros T xy2 jac_x jac_y

      ~@routine
    end
  %
```

158 In really implementation, we utilize Julia’s multiple dispatch. And “insert” the gradient code by
 159 overloading the basic instructions for the gradient wrapper type GVar. The same strategy has been
 160 used in the ForwardDiff package in Julia.

161 One does not need to define a similar function on (:+=) (log) because macro @i will generate it
 162 automatically. Notice that taking inverse and computing gradients commute [McInerney \(2015\)](#).

163 3.2 Hessians

164 Combining the uncomputing program in NiLang with dual-numbers is a simple yet efficient way to
 165 obtain Hessians. The dual number is the scalar type for computing gradients in the forward mode
 166 AD, it wraps the original scalar with a extra gradient field. The gradient field of a dual number is
 167 updated automatically as the computation marches forward. By wrapping the elementary type with
 168 Dual defined in package ForwardDiff [Revels et al. \(2016\)](#) and throwing it into the gradient program

defined in NiLang, one obtains one row/column of the Hessian matrix straightforward. We will show a benchmark in Sec. 4.2.

3.3 Complex numbers

To differentiate complex numbers, we re-implemented complex instructions reversibly. For example, with the reversible function defined in Listing. 2, we can differentiate complex valued log with no extra effort.

3.4 CUDA kernels

CUDA programming is playing a significant role in high-performance computing. In Julia, one can write GPU compatible functions in native Julia language with [KernelAbstractions Besard et al. \(2017\)](#). Since NiLang does not push variables into stack automatically for users, it is safe to write differentiable GPU kernels with NiLang. We will show this feature in the benchmarks of bundle adjustment (BA) in Sec. 4.3. Here, one should notice that the shared read in forward pass will become shared write in the backward pass, which may result in incorrect gradients. We will review this issue in the supplementary material.

4 Benchmarks

It is interesting to see how does our framework comparing with the state-of-the-art GP-AD frameworks, including source code transformation based Tapenade and Zygote and operator overloading based ForwardDiff and ReverseDiff. Since most DS-AD packages like famous Tensorflow and Pytorch are not designed for the following using cases, we do not benchmark those package. In the following benchmarks, the CPU device is Intel(R) Xeon(R) Gold 6230 CPU @ 2.10GHz, and the GPU device is Nvidia Titan V. For NiLang benchmarks, we have turned off the reversibility check off to achieve a better performance. Codes used in benchmarks could be found in the Examples section of the supplementary material.

4.1 Sparse matrices

We benchmarked the call, uncall and backward propagation time used for sparse matrix dot product and matrix multiplication. Here, we estimate the time for back propagating gradients rather than including both forward and backward, since `mul!` does not output a scalar as loss.

	dot	mul! (complex valued)
Julia-O	3.493e-04	8.005e-05
NiLang-O	4.675e-04	9.332e-05
NiLang-B	5.821e-04	2.214e-04

Table 1: Absolute runtimes in seconds for computing the objectives (O) and the backward pass (B) of sparse matrix operations. The matrix size is 1000×1000 , and the element density is 0.05. The total time used in computing gradient can be estimated by summing “O” and “B”.

The time used for computing backward pass is approximately 1.5-3 times the Julia’s native forward pass. This is because the instruction length of differentiating basic arithmetic instructions is longer than pure computing.

4.2 Graph embedding problem

Since one can combine ForwardDiff and NiLang to obtain Hessians, it is interesting to see how much performance we can get in differentiating the graph embedding program. The problem definition could be found in the supplementary material.

In Table 2, we show the performance of different implementations by varying the dimension k . The number of parameters is $10k$. As the baseline, (a) shows the time for computing the function

k	2	4	6	8	10
Julia-O	4.477e-06	4.729e-06	4.959e-06	5.196e-06	5.567e-06
NiLang-O	7.173e-06	7.783e-06	8.558e-06	9.212e-06	1.002e-05
NiLang-U	7.453e-06	7.839e-06	8.464e-06	9.298e-06	1.054e-05
NiLang-G	1.509e-05	1.690e-05	1.872e-05	2.076e-05	2.266e-05
ReverseDiff-G	2.823e-05	4.582e-05	6.045e-05	7.651e-05	9.666e-05
ForwardDiff-G	1.518e-05	4.053e-05	6.732e-05	1.184e-04	1.701e-04
Zygote-G	5.315e-04	5.570e-04	5.811e-04	6.096e-04	6.396e-04
(NiLang+F)-H	4.528e-04	1.025e-03	1.740e-03	2.577e-03	3.558e-03
ForwardDiff-H	2.378e-04	2.380e-03	6.903e-03	1.967e-02	3.978e-02
(ReverseDiff+F)-H	1.966e-03	6.058e-03	1.225e-02	2.035e-02	3.140e-02

Table 2: Absolute times in seconds for computing the objectives (O), uncall objective (U), gradients (G) and Hessians (H) of the graph embedding program. k is the embedding dimension, the number of parameters is $10k$.

call. We have reversible and irreversible implementations, where the reversible program is slower than the irreversible native Julia program by a factor of ~ 2 due to the uncomputing overhead. The reversible program shows the advantage of obtaining gradients when the dimension $k \geq 3$. The larger the number of inputs, the more advantage it shows due to the overhead proportional to input size in forward mode AD. The same reason applies to computing Hessians, where the combo of NiLang and ForwardDiff gives the best performance for $k \geq 3$.

4.3 Gaussian mixture model and bundle adjustment

We reproduced the benchmarks for Gaussian mixture model (GMM) and bundle adjustment (BA) in [Srajer et al. \(2018\)](#) by re-writing the programs in a reversible style. We show the results in Table 3 and Table 4. In our new benchmarks, we also rewrite the ForwardDiff program for a fair benchmark, this explains the difference between our results and the original benchmark. The Tapenade data is obtained by executing the docker file provided by the original benchmark, which provides a baseline for comparison.

# parameters	3.00e+1	3.30e+2	1.20e+3	3.30e+3	1.07e+4	2.15e+4	5.36e+4	4.29e+5
Julia-O	9.844e-03	1.166e-02	2.797e-01	9.745e-02	3.903e-02	7.476e-02	2.284e-01	3.593e+00
NiLang-O	3.655e-03	1.425e-02	1.040e-01	1.389e-01	7.388e-02	1.491e-01	4.176e-01	5.462e+00
Tapende-O	1.484e-03	3.747e-03	4.836e-02	3.578e-02	5.314e-02	1.069e-01	2.583e-01	2.200e+00
ForwardDiff-G	3.551e-02	1.673e+00	4.811e+01	1.599e+02	-	-	-	-
NiLang-G	9.102e-03	3.709e-02	2.830e-01	3.556e-01	6.652e-01	1.449e+00	3.590e+00	3.342e+01
Tapenade-G	5.484e-03	1.434e-02	2.205e-01	1.497e-01	4.396e-01	9.588e-01	2.586e+00	2.442e+01

Table 3: Absolute runtimes in seconds for computing the objective (O) and gradients (G) of GMM with 10k data points. “-” represents missing data due to not finishing the computing in limited time.

In the GMM benchmark, NiLang’s objective function has overhead comparing with irreversible programs in most cases. Except the uncomputing overhead, it is also because our naive reversible matrix-vector multiplication is much slower than the highly optimized BLAS function, where the matrix-vector multiplication is the bottleneck of the computation. The forward mode AD suffers from too large input dimension in the large number of parameters regime. Although ForwardDiff batches the gradient fields, the overhead proportional to input size still dominates. The source to source AD framework Tapenade is faster than NiLang in all scales of input parameters, but the ratio between computing the gradients and the objective function are close.

In the BA benchmark, reverse mode AD shows slight advantage over ForwardDiff. The bottleneck of computing this large sparse Jacobian is computing the Jacobian of a elementary function with 15 input arguments and 2 output arguments, where input space is larger than the output space. In this instance, our reversible implementation is even faster than the source code transformation based AD

# measurements	3.18e+4	2.04e+5	2.87e+5	5.64e+5	1.09e+6	4.75e+6	9.13e+6
Julia-O	2.020e-03	1.292e-02	1.812e-02	3.563e-02	6.904e-02	3.447e-01	6.671e-01
NiLang-O	2.708e-03	1.757e-02	2.438e-02	4.877e-02	9.536e-02	4.170e-01	8.020e-01
Tapenade-O	1.632e-03	1.056e-02	1.540e-02	2.927e-02	5.687e-02	2.481e-01	4.780e-01
ForwardDiff-J	6.579e-02	5.342e-01	7.369e-01	1.469e+00	2.878e+00	1.294e+01	2.648e+01
NiLang-J	1.651e-02	1.182e-01	1.668e-01	3.273e-01	6.375e-01	2.785e+00	5.535e+00
NiLang-J (GPU)	1.354e-04	4.329e-04	5.997e-04	1.735e-03	2.861e-03	1.021e-02	2.179e-02
Tapenade-J	1.940e-02	1.255e-01	1.769e-01	3.489e-01	6.720e-01	2.935e+00	6.027e+00

Table 4: Absolute runtimes in seconds for computing the objective (O) and Jacobians (J) in bundle adjustment.

framework Tapenade. With KernelAbstractions, we run our zero allocation reversible program on GPU, which provides a >200x speed up.

Broader Impact

Our automatic differentiation in a reversible eDSL brings the field of reversible computing to the modern context. We believe it will be accepted by the public to meet current scientific automatic differentiation needs and aim for future energy-efficient reversible devices. For solving practical issues, in an unpublished paper, we have successfully differentiated a spin-glass solver to find the optimal configuration on a 28×28 square lattice in a reasonable time. There are also some interesting applications like normalizing flow and bundle adjustment in the example folder of [NiLang](#) repository and [JuliaReverse](#) organization. For the future, energy consumption is an even more fundamental issue than computing time and memory. Current computing devices, including CPU, GPU, TPU, and NPU consume much energy, which will finally hit the "energy wall". We must get prepared for the technical evolution of reversible computing (quantum or classical), which may cost several orders less energy than current devices.

We also see some drawbacks to the current design. It requires the programmer to change to programing style rather than put effort into optimizing regular codes. It is not fully compatible with modern software stacks. Everything, including instruction sets and BLAS functions, should be redesigned to support reversible programming better. We put more potential issues and opportunities in the discussion section of the supplementary material. Solving these issues requires the participation of people from multiple fields.

References

- M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke, Y. Yu, and X. Zheng, “[TensorFlow: Large-scale machine learning on heterogeneous systems,](#)” (2015), software available from tensorflow.org.
- A. Paszke, S. Gross, S. Chintala, G. Chanan, E. Yang, Z. DeVito, Z. Lin, A. Desmaison, L. Antiga, and A. Lerer, in *NIPS Autodiff Workshop* (2017).
- M. Innes, E. Saba, K. Fischer, D. Gandhi, M. C. Rudilosso, N. M. Joy, T. Karmali, A. Pal, and V. Shah, “Fashionable modelling with flux,” (2018), [arXiv:1811.01457 \[cs.PL\]](#).
- J.-G. L. Hao Xie and L. Wang, [arXiv:2001.04121](#).
- G. H. Golub and C. F. Van Loan, *Matrix computations*, Vol. 3 (JHU press, 2012).
- H.-J. Liao, J.-G. Liu, L. Wang, and T. Xiang, *Physical Review X* **9** (2019), [10.1103/physrevx.9.031041](#).

266 M. Seeger, A. Hetzel, Z. Dai, E. Meissner, and N. D. Lawrence, “Auto-differentiating linear
267 algebra,” (2017), [arXiv:1710.08717 \[cs.MS\]](#) .

268 Z.-Q. Wan and S.-X. Zhang, “Automatic differentiation for complex valued svd,” (2019),
269 [arXiv:1909.02659 \[math.NA\]](#) .

270 C. Hubig, “Use and implementation of autodifferentiation in tensor network methods with complex
271 scalars,” (2019), [arXiv:1907.13422 \[cond-mat.str-el\]](#) .

272 X.-Z. Luo, J.-G. Liu, P. Zhang, and L. Wang, “Yao.jl: Extensible, efficient framework for quantum
273 algorithm design,” (2019), [arXiv:1912.10877 \[quant-ph\]](#) .

274 L. Hascoet and V. Pascual, *ACM Transactions on Mathematical Software (TOMS)* **39**, 20 (2013).

275 M. Innes, “Don’t unroll adjoint: Differentiating ssa-form programs,” (2018), [arXiv:1810.07951](#)
276 [\[cs.PL\]](#) .

277 M. Innes, A. Edelman, K. Fischer, C. Rackauckas, E. Saba, V. B. Shah, and W. Tebbutt, *CoRR*
278 [abs/1907.07587](#) (2019), [arXiv:1907.07587](#) .

279 J. Bezanson, S. Karpinski, V. B. Shah, and A. Edelman, *arXiv preprint arXiv:1209.5145* (2012).

280 J. Bezanson, A. Edelman, S. Karpinski, and V. B. Shah, *SIAM Review* **59**, 65–98 (2017).

281 K. S. Perumalla, *Introduction to reversible computing* (Chapman and Hall/CRC, 2013).

282 M. P. Frank, *IEEE Spectrum* **54**, 32–37 (2017).

283 C. Lutz, “Janus: a time-reversible language,” (1986), *Letter to R. Landauer*.

284 M. P. Frank, *The R programming language and compiler*, Tech. Rep. (MIT Reversible Computing
285 Project Memo, 1997).

286 I. Lanese, N. Nishida, A. Palacios, and G. Vidal, *Journal of Logical and Algebraic Methods in*
287 *Programming* **100**, 71–97 (2018).

288 T. Haulund, “Design and implementation of a reversible object-oriented programming language,”
289 (2017), [arXiv:1707.07845 \[cs.PL\]](#) .

290 M. MacKay, P. Vicol, J. Ba, and R. B. Grosse, in *Advances in Neural Information Processing*
291 *Systems 31*, edited by S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and
292 R. Garnett (Curran Associates, Inc., 2018) pp. 9029–9040.

293 L. Dinh, D. Krueger, and Y. Bengio, “Nice: Non-linear independent components estimation,”
294 (2014), [arXiv:1410.8516 \[cs.LG\]](#) .

295 D. Maclaurin, D. Duvenaud, and R. Adams, in *Proceedings of the 32nd International Conference*
296 *on Machine Learning*, Proceedings of Machine Learning Research, Vol. 37, edited by F. Bach and
297 D. Blei (PMLR, Lille, France, 2015) pp. 2113–2122.

298 J. Behrmann, D. Duvenaud, and J. Jacobsen, *CoRR abs/1811.00995* (2018), [arXiv:1811.00995](#) .

299 M. P. Frank and T. F. Knight Jr, *Reversibility for efficient computing*, Ph.D. thesis, Massachusetts
300 Institute of Technology, Dept. of Electrical Engineering and ... (1999).

301 J. G. Koller and W. C. Athas, in *Workshop on Physics and Computation* (1992) pp. 267–270.

302 R. C. Merkle, R. A. Freitas, T. Hogg, T. E. Moore, M. S. Moses, and J. Ryley, *Journal of*
303 *Mechanisms and Robotics* **10** (2018), 10.1115/1.4041209.

304 K. Likharev, *IEEE Transactions on Magnetics* **13**, 242 (1977).

305 V. K. Semenov, G. V. Danilov, and D. V. Averin, *IEEE Transactions on Applied Superconductivity*
306 **13**, 938 (2003).

307 R. Landauer, *IBM journal of research and development* **5**, 183 (1961).

- 308 C. H. Bennett (1973).
- 309 D. Speyer and B. Sturmfels, [Mathematics Magazine](#) **82**, 163 (2009).
- 310 A. McInerney, *First steps in differential geometry* (Springer, 2015).
- 311 J. Revels, M. Lubin, and T. Papamarkou, “Forward-mode automatic differentiation in julia,” (2016),
312 [arXiv:1607.07892 \[cs.MS\]](#) .
- 313 T. Besard, C. Foket, and B. D. Sutter, [CoRR abs/1712.03112](#) (2017), [arXiv:1712.03112](#) .
- 314 F. Srajer, Z. Kukelova, and A. Fitzgibbon, *Optimization Methods and Software* **33**, 889 (2018).