Instruction level automatic differentiation on a reversible Turing machine

Jin-Guo Liu,[1, *] Taine Zhao,[2] and [**JG: MLS's name here**][3]

[1]*Institute of Physics, Chinese Academy of Sciences, Beijing 100190, China*
[2]*Department of Computer Science, University of Tsukuba*
[3]*MLS's affiliations here*

This paper considers source to source automatic differentiation (AD) on a reversible Turing machine. We start by reviewing why adjoint mode AD is hard for traditional machine learning frameworks. Then we propose a solution to existing issues by writing a program reversiblly. We developed an reversible eDSL NiLang in Julia that can differentiate a general program while being compatible with Julia ecosystem. It can generate efficient backward rules with reasonable time and memory efficiency. We demonstrate its power by generating the backard rules for Bessel function, sparse matrix dot product, unitary matrix multiplication and QR decomposition. It can also be used to solve a graph embeding problem efficiently. Also, we will discuss the challenges that we face towards rigorous reversible programming from the instruction and hardware perspective.

I. INTRODUCTION

Computing the gradients of a numeric model $f : \mathbb{R}^m \to \mathbb{R}^n$ plays a crucial role in scientific computing. Consider a computing process

$$\mathbf{x}^1 = f_1(\mathbf{x}^0)$$
$$\mathbf{x}^2 = f_2(\mathbf{x}^1)$$
$$\cdots$$
$$\mathbf{x}^L = f_L(\mathbf{x}^{L-1})$$

where $x^0 \in R^m$, $x^L \in R^n$, L is the depth of computing. The Jacobian of this program is a $n \times m$ matrix $J_{ij} \equiv \frac{\partial x_i^L}{\partial x_j^0}$, where x_j^0 and x_i^L are single elements of inputs and outputs. Computing part of the Jacobian automatically is what we called automatic differentiation (AD). It can be classified into three classes, the tangent mode AD, the adjoint mode AD and the mixed mode AD. [1] The tangent mode AD computes the Jacobian matrix elements that related to a single input using the chain rule $\frac{\partial \mathbf{x}^k}{\partial x_j^0} = \frac{\partial \mathbf{x}^k}{\partial \mathbf{x}^{k-1}} \frac{\partial \mathbf{x}^{k-1}}{\partial x_j^0}$, while a tangent mode AD computes Jacobian matrix elements that related to a single output using $\frac{\partial \mathbf{x}^k}{\partial x_j^0} = \frac{\partial \mathbf{x}^k}{\partial \mathbf{x}^{k-1}} \frac{\partial \mathbf{x}^{k-1}}{\partial x_j^0}$. In variational applications where the loss function always outputs a scalar, the adjoint mode AD is preferred. However, implementing adjoint mode AD is harder than implementing its tangent mode counterpart, because it requires propagating the gradients in the inverse direction of computing the loss. The back propagation of gradients requires intermediate information of a program that includes

1. the computational process,

2. and variables used in computing gradients.

The computational process is often stored in a computational graph, a directed acyclic graph (DAG) that represents the relationship between data and functions. In Pytorch [2] and Flux [3], every variable has a tracker field that stores its parent information, i.e., the input data and function generating this variable. TensorFlow [4] implements a static computational graph as a description of the program before actual computation happens. The required variables are also recorded in this graph. For source to source AD package, Tools like Tapenade [1] and Zygote [5, 6] generate the adjoint code statically while putting variable on a stack called the wengert list. [**JG: The introduction to source to source AD is a bit weak, please give reader an overview about the development and difficulty in this field.**] As reverse mode AD require storing intermediate result, the memory consumption of the program is typically O(n) where n is the runtime of the program. In many deep learning models like recurrent neural network [7] and residual neural networks [8], the depth can reach several thousand, where the memory is often the bottleneck of these programs. In traditional AD applications, a program that might do billions of computations will as well get a wengert list in the range of GBs, thus memory will become a bottleneck as well. The memory management is becoming a wall [9] that limiting the scale of many applications on GPU. The computational power of a Nvidia V100 GPU can reaches 100 TFLOPS, competing a small cluster. However, its memory is only 32GB. A better solution to memory management should be found. [**JG: @MLS, Do you have any good comments about starving CPU, i.e. memory accessing is consuming much more time than computing itself. This is why caching intermediate result is a performance killer.**]

We tackle this problem by making a program reversible. In the machine learning field, reversibility has been used in reduce the memory allocations in recurrent neural networks /[10] and residual neural networks [11]. These techinics include information buffer [12] and reversible activation functions [13, 14]. However, our approach is general purposed. We develop an embedded domain-specific language (eDSL) NiLang in Julia language [15, 16] that implements reversible programming. [17, 18]. This eDSL provides a macro to generate reversible functions, one can write reversible control flows, instructions and memory managements in this macro. In history, there have been some prototypes of reversible languages like Janus [19], R (not the popular one) [20], Erlang [21] and object-oriented ROOPL [22]. In

the past, the primary motivation of studying reversible programming is to support reversible devices [23] like adiabatic complementary metal–oxide–semiconductor (CMOS) [24], molecular mechanical computing system [25] and superconducting system [26, 27]. These devices either implements reversible logical gates or is able to recover signal energy, where the latter is also called generalized reversible computing. Reversible computing are more energy efficient from information and entropy's perspective by the Landauer's principle [28] After decades of efforts, reversible computing devices are very close to providing productivity now. For exmaple, adiabatic CMOS can be a better choice already in a spacecraft [29], where energy is more valuable than device itself. Reversible programming are also interesting to software engineers, it is a powerful tool to schedule asynchronious events [30] and debug a program bidirectionally [31]. However, the field of reversible computing faces the difficulty of not enough funding recent decade [18]. As a result, little people know what this field did in machine learning field. People haven't connected it with automatic differentiation seriously even though they are so related. We aims to breaks the information barrier between the machine learning community and the reversible programming community in our work, and provides yet another strong motivation to develop reversible programming.

In this paper, we first introduce the language design of NiLang in Sec. II. In Sec. III, we explain the back-propagation algorithm of Jacobians and Hessians in this eDSL. In Sec. IV, we show several examples including Fibonacci number, Bessel function, unitary matrix multiplication and QR decomposition [32]. We show how to generate first order and second order backward rules for these functions. In Sec. V, we discuss several important issues, how time-space tradeoff works, reversible instructions and hardware, and finally, an outlook to some open problems to be solved. In the appendix, we show the grammar of NiLang and a gradient free self-consistent training strategy.

II. LANGUAGE DESIGN

A. Introductions to reversible language design

[**JG: how about now? feel free to comment more**] [**MLS: Add more introduction to reversible programming**] In a modern programming language, functions are pushed to a global stack for scheduling. The memory layout of a function consists of input arguments, a function frame with information like the return address and saved memory segments, local variables, and working stack. After the call, the function clears these run-time information, only stores the return value. In the reversible programming, this kind of design is no longer the best practice. One can not discard input variables and local variables easily after a function call, since discarding information may ruin reversibility. For this reason, reversible functions are very different from irreversible ones from multiple perspectives.
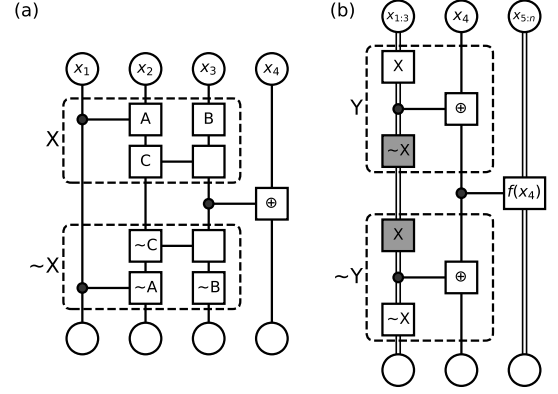

Figure 1: Data oriented computational graph that shows the compute-copy-uncompute paradigm. Circles connected by vertical lines are variables, where double lines represents multiple variables. Boxes and dots connected by parallel lines are operations. A dot on a line represents a control parameter while a box represents a mutable parameter.

1. Memory management

The key difference between reversible and irreversible memory management is the content of a variable that going to be discarded in a reversible program must be known. We denote the allocation of a zero emptied memory as $x \leftarrow 0$, and the corresponding deallocation as $x \rightarrow 0$. A variable x can be allocated and deallocated in a local scope, which is called an ancilla. It can also be push to a stack and used later with a pop statement. The reversible stack is similar to a traditional stack except it zero-clears the variable after pushing and presupposes the variable being zero-cleared before popping.

Knowing the contents in the memory when deallocating is not easy, hence Charles H. Bennett introduced the famous compute-copy-uncompute paradigm [33]. In order to show how reversible memory manage works better, we introduce the memory oriented computational graph as shown in Fig. 1. Notations are highly inspired by quantum circuit representations. A vertical line is a variable, it can be used by multiple operations, which means it is a hypergraph rather than a simple graph like DAG. When a variable is used by a function, we differentiate parameters by whether its value is changed by this function. We put a box if it is changed, other wise a dot. Panel (a) shows that in the computing stage, the program in dashed box X is executed on space $x_{1:3}$ to compute desired result. In the copying stage, the contents in x_3 is read out to a pre-emptied memory x_4 through addition operation $\oplus$, because this is the piece of information that we care. Since this copy operation does not change contents of $x_{1:3}$, we can use the inverse operation ~X to undo all the changes to these registers. If some of these uncomputed variables are initialized to some known value like 0, now they can be deallocated since their states is known again. If this subroutine of generating x_4 is used in another program as
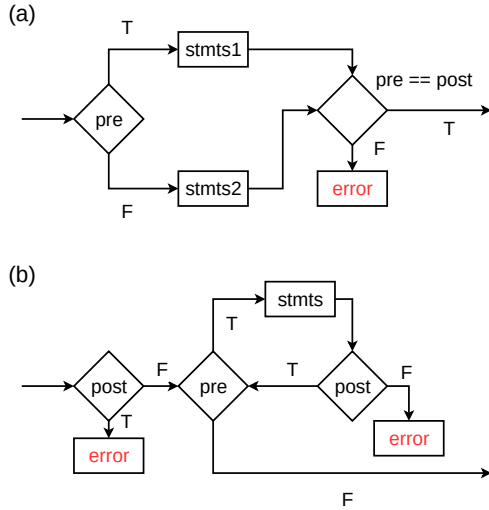
Figure 2: Flow chart for reversible (a) `if` statement and (b)
`while` statement. "stmts", "stmts1" and "stmts2" are
statements, statements in true branch and statements in false
branch respectively. "pre" and "post" are precondition and
postconditions respectively.

3. Arithmetic instructions

Every instruction has a unique inverse that can undo the
changes. For logical operations, we have `y ⊻= f(args...)`
self reversible. For other arithmetic operations, we assume
`y += f(args...)` and `y -= f(args...)` are reversible to
each other although it is not true for floating point numbers
due to the rounding error. Here `f` can be `identity`, `*`, `/` and
`^` et. al. Besides the above two types of operations, `SWAP`
operation that exchanges the contents in two memory spaces
are also widely used in reversible computing systems. We
will have a detailed discussion about the rounding error and
number system later in Sec. V B

shown in Fig. 1 (b), x_4 can be uncomputed by the reversing
the whole subroutine in panel (a). The interesting fact is, both
X and $\sim X$ are executed twice, which is not necessary. We can
of cause cancel a pair of X and $\sim X$ (the gray boxes), by doing
this, we have to keep the memory $x_{1:3}$ during computing
$f(x_4)$, i.e. additional space. The trade off between space and
time will discuss the detail in Sec. V A.

B. NiLang's Reversible IR

In the last subsection, we have reviewed basic building
blocks of a typical reversible language. In order to insert
the code of obtaining gradients into the reversed program,
the reversible language design should have related abstraction
power. This motivates us to design a new reversible language
NiLang to fit this task. NiLang is an eDSL in Julia. We choose
Julia as the host language for multiple purposes. Julia's
meta programming and its package for pattern matching
MLStyle.jl [34] allow us to define an eDSL conveniently.
Meanwhile, the type inference and just in time compiling can
remove most overheads introduced in our eDSL, providing
a reasonable performance. Most importantly, the multiple
dispatch provides the polymorphism that will be used in our
autodiff engine.

The main feature of NiLang is contained in a single macro
`@i` that compiles a reversible function. The allowed state-
ments in this eDSL are shown in Appendix A. The following
is a minimal example of compiling a NiLang function to
native julia function.

2. Control flows

The reversible `if` statement is shown in Fig. 2 (a). It
contains two conditions, a precondition and a postcondition.
The precondition decides which branch to enter in the forward
execution, while the postcondition decides which branch
to enter in the backward execution. After executing the
specific branch, the program checks the consistency between
precondition and postcondition to make sure they are the
consistent. The reversible `while` statement is shown in Fig. 2
(b). It also has both precondition and postcondition. Before
executing the condition expressions, the program preassumes
the postcondition is false. After each iteration, the program
asserts the postcondition to be true. In the reverse pass, we
exchange the precondition and postcondition. The reversible
`for` statement is similar to irreversible ones except after
executing the loop, the program checks the values of these
variables to make sure they are not changed. In the reverse
pass, we exchange start and stop and inverse the sign of step.

```
julia> using NiLangCore, MacroTools

julia> macroexpand(Main, :(@i function f(x, y)
           SWAP(x, y)
       end)) |> MacroTools.prettify
quote
    $(Expr(:meta, :doc))
    function $(Expr(:where, :(f(x, y))))
        gaur = SWAP(x, y)
        x = (NiLangCore.wrap_tuple(gaur))[1]
        y = (NiLangCore.wrap_tuple(gaur))[2]
        return (x, y)
    end
    if typeof(f) != typeof(~f)

     function $(Expr(:where, :((  #$ TODO: remove this comment
                    mongoose::typeof(~f))(x, y))))
            mandrill = (~SWAP)(x, y)
            x = (NiLangCore.wrap_tuple(mandrill))[1]
            y = (NiLangCore.wrap_tuple(mandrill))[2]
            return (x, y)
        end
    end
    if !(NiLangCore._hasmethod1(
              NiLangCore.isreversible, typeof(f)))
        NiLangCore.isreversible(::typeof(f)) = true
    end
end
```

statement	inverse
`<f>(<args>...)`	`(~<f>)(<args>...)`
`<y> += <f>(<args>...)`	`<y> -= <f>(<args>...)`
`<y> .+= <f>.(<args>...)`	`<y> .-= <f>.(<args>...)`
`<y> ⊻= <f>(<args>...)`	`<y> ⊻= <f>(<args>...)`
`<y> .⊻= <f>(<args>...)`	`<y> .⊻= <f>(<args>...)`
`<a> ← <expr>`	`<a> → <expr>`
`(<T1> => <T2>)(<x>)`	`(<T2> => <T1>)(<x>)`

`begin`
`  <stmts>`
`end`
	`begin`
	`  ~(<stmts>)`
	`end`

`if (<pre>, <post>)`
`  <stmts1>`
`else`
`  <stmts2>`
`end`
	`if (<post>, <pre>)`
	`  ~(<stmts1>)`
	`else`
	`  ~(<stmts2>)`
	`end`

`while (<pre>, <post>)`
`  <stmts>`
`end`
	`while (<post>, <pre>)`
	`  ~(<stmts>)`
	`end`

`for <i>=<m>:<s>:<n>`
`  <stmts>`
`end`
	`for <i>=<m>:-<s>:<n>`
	`  ~(<stmts>)`
	`end`

`@safe <expr>`	`@safe <expr>`

Table I: A collection of reversible statements. "." is the symbol for broadcasting magic in Julia, "~" is the symbol for reversing a statement or a function. `<...>` represents a non-keyword, where `<pre>` stands for precondition, `<post>` stands for postcondition, `<args>...` stands for the argument list of a function, `<stmts>` stands for statement, `<exprs>` stands for expression, `<T1>` and `<T2>` stand for types and the reset are variables.

,

Macro `@i` generates three functions `f`, `~f` and `NiLangCore.isreversible`. `f` and `~f` are a pair of functions that reverse to each other, where `~f` is an callable of type `Inv{typeof(f)}`. In the body of `f`, `NiLangCore.wrap_tuple` is used to unify output data types, it will wrap any non-tuple variable to a tuple. The outputs of `SWAP` are assigned back to its input variables, in other words, a function modifies inputs inplace. At the end this function, this macro attaches a return statement that returns all input variables. `NiLangCore.isreversible` is a function to mark the reversibility trait of `f`.

To understand the design of reversibility, we first introduce a reversible IR that plays a central role in NiLang. In this IR, a statement can be an instruction, a function call, a control flow, a memory allocation/deallocation, or the inverse statement "~". Any statement is this IR has a unique inverse as shown in Table I.

"←" and "→" are symbols for memory allocation and deallocation, one can input them by typing "\leftarrow" and "\rightarrow" respectively followed by a Tab key in a Julia editor or REPL. "`begin <stmts> end`" is the block statement in Julia, it represents a code block. It can be inverted by reversing the order of `<stmts>` as well as each element in it. The conditional expression in `if` or `while` statements is a tuple of precondition and postcondition. Finally, the special macro `@safe` allows users to use external statements that do not break reversibility. For example, one can use `@safe @show <var>` for debugging.

C. Compiling

The compilation of a reversible function contains three stages.

The first stage pre-processes human inputs to a reversible IR. The preprocessor expands the symbol "~" in the postcondition field of `if` statement by copying the precondition, adds missing ancilla "←" statements to ensure "←" and "→" appear in pairs inside a function, a `while` statement or a `for` statement, and expands the uncomputing macro `~@routine`. Since the compmute-copy-uncompute paradigm is extensively used in reversible programming for uncomputing ancillas. One can use `@routine <stmt>` statement to record a statement, and `~@routine` to insert `~<stmt>` for uncomputing. The following example pre-processes an `if` statement to the reversible IR.

```
julia> using NiLangCore, MacroTools

julia> MacroTools.prettify(
           @code_preprocess if (x > 3, ~)
               @routine z += x * y
               ~@routine
           end)
:(if (x > 3, x > 3)
      z += x * y
      z -= x * y
  else
  end)
```

```
julia> MacroTools.prettify(
           @code_interpret if (pre, post)
               z += x * y
           else
               z += x / y
           end)
quote
    bat = pre
    if bat
        @assignback (PlusEq(*))(z, x, y)
    else
        @assignback (PlusEq(/))(z, x, y)
    end
    @invcheck post bat
end
```

In this example, since the precondition "x > 3" is not change after execution of the specific branch, we omit the postcondition by putting a "~" in this field. "@routine" records a statement, the statement can also be a "begin <stmts> end" block as a sequence of statements.

The second stage generates the reversed code according to table Table I. For example,

The compiler translates the instruction according to Table III and adds @assignback before each instruction and function call statement. The macro @assignback assigns the output of a function back to the arguments of that function. @invcheck post bat checks the consistency between preconditions and postconditions to ensure reversibility. This statement will throw an InvertibilityError error if target variables bat and post are not "equal" to each other up to a certian torlerance.

D. Types and Dataviews

So far, the language design is not too different from a traditional reversible language. To implement the adjoint mode AD, we introduce types and dataviews. The type that used in the reversible context is just a normal Julia type with an extra requirement of having reversible constructors. The inverse of a constructor is called a "destructor", which unpacks data and deallocates derived fields. Data packing is implemented by reinterpreting the new function in Julia. For example,

```
julia> MacroTools.prettify(
           @code_reverse if (pre, post)
               z += x * y
           else
               z += x / y
           end)
:(if (post, pre)
      z -= x * y
  else
      z -= x / y
  end)
```

```
x ← new{TX, TG}(x, g)
```

Here, the "←" statement followed by a new function is treated specially that it deallocates g. This makes sense because the output of new keeps all information in input argument list. Its inverse is

```
x → new{TX, TG}(x, g)
```

The third stage is translating this IR and its inverse to native Julia code. It explains all functions as inplace and insert codes about reversibility check. At the end of a function definition, it attaches a return statement that returns all input arguments. After this, the function is ready to execute on the host language. The following example shows how an if statement is transformed in this stage.

It unpacks x and allocates a new ancilla g. The following example shows how to define as reversible type GVar.

```
julia> using NiLangCore

julia> @i struct GVar{T,GT} <: IWrapper{T}
           x::T
           g::GT
           function GVar{T,GT}(x::T, g::GT) where
                                            {T,GT}
               new{T,GT}(x, g)
           end
           function GVar(x::T, g::GT) where {T,GT}
               new{T,GT}(x, g)
           end
           @i function GVar(x::T) where T
               g ← zero(x)
               x ← new{T,T}(x, g)
           end
           @i function GVar(x::AbstractArray)
               GVar.(x)
           end
       end

julia> GVar(0.5)
GVar{Float64,Float64}(0.5, 0.0)

julia> (~GVar)(GVar(0.5))
0.5

julia> (~GVar)(GVar([0.5, 0.6]))
2-element Array{Float64,1}:
 0.5
 0.6
```

This piece of code is copied from the autodiff submodule of NiLang. `GVar` is a type used to store gradient information of a variable. Here, we put `@i` macro before both `struct` and `function` statements. The ones before functions mark reversible functions, while the one before `struct` keyword is used to handle the function scope issue. It moves `~GVar` functions outside of this type definition.

The reversible cast between two types can be defined conveniently with the macro `@icast`.

```
julia> @pure_wrapper A

julia> @icast A(x) => GVar(x, g) begin
           g ← zero(x)
           g += identity(1)
       end

julia> x = A(0.5)
A(0.5)

julia> @instr (A=>GVar)(x)
GVar{Float64,Float64}(0.5, 1.0)

julia> @instr (GVar=>A)(x)
A(0.5)
```

Here, we first define a simple reversible wrapper `A` using

macro `@pure_wrapper`, and then the cast rule between `A` type and `GVar` type. The body of cast is a reversible mapping that transforms `x` to `(x, g)`. The compiler appends a default constructor `DVar(xx, gg)` at the end of program to instantiate a new object as the return value. Its inverse that coverts an object of type `GVar` to type `A` is automatically generated by reversing the above statements.

The fields of an object can be accessed and altered through dataviews. A dataview can be an object, a field of a dataview, an array element of a dataview, or a bijective mapping of a dataview. Let us first consider the following example.

```
julia> arr = [GVar(3.0), GVar(1.0)]
2-element Array{GVar{Float64,Float64},1}:
 GVar{Float64,Float64}(3.0, 0.0)
 GVar{Float64,Float64}(1.0, 0.0)

julia> x, y = 1.0, 2.0
(1.0, 2.0)

julia> @instr -arr[2].g += x * y
2.0

julia> arr
2-element Array{GVar{Float64,Float64},1}:
 GVar{Float64,Float64}(3.0, 0.0)
 GVar{Float64,Float64}(1.0, -2.0)
```

Here, both `-arr[2].g`, `x` and `y` are dataviews. In Julia language, the statement `-grad(arr[2]) += x * y` should throw a syntax error because the function call "`-`" can not be assigned, and `GVar` is an immutable type. In our eDSL, we wish it works because a memory cell is assumed to be modifiable in our eDSL. The secret of how it works lies in the macro `@assignback`, it translates the above statement to

```
1  res = (PlusEq(*))(-arr[2].g, x, y)
2  arr[2] = chfield(arr[2], Val(:g),
3      chfield(arr[2].g, -, res[1]))
4  x = res[2]
5  y = res[3]
```

The first line `PlusEq(*)(-arr[3].g, x, y)` computes the output, which is a tuple of length 3. At lines 2-3, `chfield(x, Val{:g}, val)` modifies the `g` field of `x` and `chfield(x, -, res[1])` returns `-res[1]`. Here, modifying a field requires the default constructor of a type not overwritten. The assignments in lines 4 and 5 are straightforward.

III. AUTOMATIC DIFFERENTIATION

Local Jacobians and Hessians for basic instructions used in this section could be found in Appendix B 1.

Figure 3: Computational processes in the tensor network diagram, a big circle with three legs represents a Hessian, a small circle with two legs represents a Jacobian. Dangling edges and connected edges stands for unpaired and paired labels respectively in the Einstein's notation. From left to right, the diagrams represent computing, uncomputing, back propagating Jacobians and back propagating Hessians.

A. First order gradient

Consider a computation $\mathbf{x}^{i-1} = f_i^{-1}(\mathbf{x}^i)$ in a reversed program, the Jacobians can be propagated in the reversed direction like [**MLS: adjoint mode AD does not propagate Jacobians.**][**JG: now?**]

$$
\begin{aligned}
J_{\mathbf{x}^{L'}}^{\mathbf{x}^L} &= \delta_{\mathbf{x}^L, \mathbf{x}^{L'}}, \\
J_{\mathbf{x}^{i-1}}^{\mathbf{x}^L} &= J_{\mathbf{x}^i}^{\mathbf{x}^L} J_{\mathbf{x}^{i-1}}^{\mathbf{x}^i},
\end{aligned}
\tag{1}
$$

where $\mathbf{x}^L$ represents the outputs of the program, $J_{\mathbf{x}^i}^{\mathbf{x}^L} \equiv \frac{\partial \mathbf{x}^L}{\partial \mathbf{x}^i}$ is the Jacobian to be propagated, and $J_{\mathbf{x}^{i-1}}^{\mathbf{x}^i}$ is the local Jacobian matrix. The Einstein's notation [35] is used here so that the duplicated index $\mathbf{x}^i$ is summed over. Eq. (1) can be rewritten in the diagram of tensor networks [36] as shown in Fig. 3.

The algorithm to compute the adjoint mode AD can be summarized as follows.

Algorithm 1: Reversible programming AD

Result: grad.$(\mathbf{x}_g)$
let `iloss` be the index of loss variable in $\mathbf{x}$
$\mathbf{y} = f(\mathbf{x})$
$\mathbf{y}_g = $ GVar.$(\mathbf{y})$
grad$(\mathbf{y}_g[\text{iloss}])$ += 1.0
$\mathbf{x}_g = f^{-1}(\mathbf{y}_g)$

The program first computes the forward pass, and then wrap each output variable with GVar. The constructor GVar attaches a zero gradient field to a variable. If an input variable is an array, GVar will be broadcasted to each array element automatically. The line grad($\mathbf{y}_g$[iloss]) += 1.0 adds one to the gradient field of loss to initialize a single row of Jacobian as in the first line of Eq. (1). Finally, execute the inverse program f^{-1}, the gradients are stored in the grad dataview of output variables. The computation of gradients are implemented with multiple dispatch, that is, when an instruction has a GVar type in its argument list, it calls a different routine. The same trick is used in the dual number implementation of tangent mode AD [37]. In the following,

we exemplify the design by binding the backward rules to instructions $\oplus(*)$ and $\ominus(*)$

```
@i function ⊖(*)(out!::GVar, x::GVar, y::GVar)
    value(out!) -= value(x) * value(y)
    grad(x) += grad(out!) * value(y)
    grad(y) += value(x) * grad(out!)
end
```

Here, we adopt a convention that only variables ended with ! will be changed after the function call. Although the backward rule is defined on $\ominus(*)$, the compiler generates the backward rules on $\oplus(*)$ too. This reflects the fact that taking inverse and computing gradients commute to each other [38]. Hence for a general reversible function f, one can bind backward rule on either f or its inverse f^{-1}. We can check the correctness of our definition like follows.

```
julia> using NiLang, NiLang.AD

julia> a, b, y = GVar(0.5), GVar(0.6), GVar(0.9)
(GVar(0.5, 0.0), GVar(0.6, 0.0), GVar(0.9, 0.0))

julia> @instr grad(y) += identity(1.0)

julia> @instr y += a * b
GVar(0.6, -0.5)

julia> a, b, y
(GVar(0.5, -0.6), GVar(0.6, -0.5), GVar(1.2, 1.0))

julia> @instr y -= a * b
GVar(0.6, 0.0)

julia> a, b, y
(GVar(0.5, 0.0), GVar(0.6, 0.0), GVar(0.899999, 1.0))
```

Here, $J(\ominus(*)) = J(\oplus(*))^{-1}$, hence consecutively applying them restores gradient fields of variables. The implementation of Algorithm 1 is so short that we present the function definition as follows.

```
@i function (g::Grad)(args...; kwargs...)
    @safe @assert count(x -> x isa Loss, args) == 1
    iloss ← 0
    @routine for i=1:length(args)
        if (tget(args,i) isa Loss, iloss==i)
            iloss += identity(i)
            (~Loss)(tget(args,i))
        end
    end

    g.f(args...; kwargs...)
    GVar.(args)
    grad(tget(args,iloss)) += identity(1.0)
    (~g.f)(args...; kwargs...)

    ~@routine
end
```

Figure 4: The time to obtain gradient as function of loop size. $\times n$ in legend represents a rescaling of time.

The program first checks variables contain exactly one `Loss`, where `Loss` is a reversible wrapper used to mark the loss variable. Then we locates the loss variable as `iloss` and use `~Loss` unwraps the loss variable. After computing the forward pass and backward pass, `~@routine` uncomputes the ancilla `iloss` and returns the location information to the loss variable. Here, `tget(args, i)` returns the i-th element of a tuple. We forbid tuple indexing deliberately in order to avoid possible ambiguity in supporting array indexing.

The overhead of using `GVar` type is negligible thanks to Julia's multiple-dispatch and type inference. Let us consider a simple example that accumulates 1.0 to a target variable *x* for *n* times.

[**JG: Grammarly here!**]

```
julia> using NiLang, NiLang.AD, BenchmarkTools

julia> @i function prog(x, one, n::Int)
           for i=1:n
               x += identity(one)
           end
       end

julia> @benchmark prog'(Loss(0.0), 1.0, 10000)
BenchmarkTools.Trial:
  memory estimate:  1.05 KiB
  allocs estimate:  39
  --------------
  minimum time:     35.838 μs (0.00% GC)
  median time:      36.055 μs (0.00% GC)
  mean time:        36.483 μs (0.00% GC)
  maximum time:     185.973 μs (0.00% GC)
  --------------
  samples:          10000
  evals/sample:     1
```

We implement the same function with TensorFlow, PyTorch and Zygote for comparison. The code could be found in our paper's GitHub repository [39]. Benchmark results on CPU Intel(R) Xeon(R) CPU E5-2680 v4 @ 2.40GHz are

shown in Fig. 4. One can see that the NiLang implementation is unreasonably fast, it is approximately two times the forward pass written in native Julia code. Reversible programming is not always as fast as its irreversible counterparts. In practical applications, a reversible program may have memory or computation overhead. We will discuss the details of time and space trade off in Sec. V A.

B. Second-order gradient

Second-order gradients can be obtained in two different approaches.

1. Back propagating first-order gradients

Back propagating the first-order gradients is the most widely used approach to obtain the second-order gradients. Suppose the function space is closed under gradient operation, one can obtain higher-order gradients by recursively differentiating lower order gradient functions without defining new backward rules.

Figure 5: Data flow in obtaining the second-order gradient with the recursive differentiation approach. Annotations on lines are data types used in the computation.

Fig. 5 shows the data flow in the four passes of computing

Hessian. The first two passes obtains the gradients. Before entering the third pass, the program wraps each field in `GVar` with another layer of `GVar`. Then we pick a variable x_i and add 1 to the gradient field of its gradient `grad(grad(x_i))` in order to compute the i-th row of Hessian. Before entering the final pass, the ~`GVar` is called. We can not unwrap `GVar` directly because although the values of gradients have been uncomputed to zero, the gradient fields of gradients may be nonzero. Instead, we use `Partial{:x}(obj)` to take field `x` of an object without erasing memory. By repeating the above procedure for different x_i, one can obtains the full Hessian matrix.

2. Hessian propagation

A probably more efficient approach is back-propagating Hessians directly [40] using the relation

$$
\begin{aligned}
H^{\mathbf{x}^L}_{\mathbf{x}^{L'},\mathbf{x}^{L''}} &= \mathbf{0}, \\
H^{\mathbf{x}^L}_{\mathbf{x}^{i-1},\mathbf{x}^{i-1'}} &= J^{\mathbf{x}^i}_{\mathbf{x}^{i-1}} H^{\mathbf{x}^L}_{\mathbf{x}^i,\mathbf{x}^{i'}} J^{\mathbf{x}^{i'}}_{\mathbf{x}^{i-1'}} + J^{\mathbf{x}^L}_{\mathbf{x}^i} H^{\mathbf{x}^i}_{\mathbf{x}^{i-1},\mathbf{x}^{i-1'}}.
\end{aligned}
\tag{2}
$$

Here, the Hessian tensor $H^{\mathbf{x}^L}_{\mathbf{x}^{i-1},\mathbf{x}^{i-1'}}$ is rank three, where the top index is often taken as a scalar and omitted. In tensor network diagram, the above equation can be represented as the right panel of Fig. 3. Hessian propagation is a special case of Taylor propagation. With respect to the order of gradients, Taylor propagation is exponentially more efficient in obtaining higher-order gradients than differentiating lower order gradients recursively. Comparing with operator overloading, source to source automatic differentiation has the advantage of having very limited primitives, exhausted implementation of Hessian propagation is possible. An example to obtain Hessians is provided in Sec. IV A.

C. Gradient on ancilla problem

In this section, we introduced an easily overlooked problem in our reversible AD framework. An ancilla sometimes can carry a nonzero gradient when it is going to be deallocated. As a result, even if an ancilla can be uncomputed rigorously in the original program, its `GVar` wrapped version is not necessarily safely deallocated. In NiLang, we simply "drop" the gradient field of ancillas instead of raising an error. In the following, we justify our approach by proving the following theorem

Theorem 1. *Deallocating an ancilla with emptied value field and nonzero gradient field does not harm the reversibility of a function.*

Proof. Consider a reversible function $\mathbf{x}^i, b = f_i(\mathbf{x}^{i-1}, a)$, where a and b are the input and output values of an ancilla. Since the ancilla is emptied for any input $\mathbf{x}^{i-1}$, we have

$$
\frac{\partial b}{\partial \mathbf{x}^{i-1}} = \mathbf{0}. \tag{3}
$$

Since the gradient fields are derived from the value fields of variables, discarding gradients should not have any effect to the `value` fields. The rest is to show `grad(b)` $\equiv \frac{\partial \mathbf{x}^L}{\partial b}$ does appear in the backward rule of this function. It can be seen from the back-propagation rule

$$
\frac{\partial \mathbf{x}^L}{\partial \mathbf{x}^{i-1}} = \frac{\partial \mathbf{x}^L}{\partial \mathbf{x}^i}\frac{\partial \mathbf{x}^i}{\partial \mathbf{x}^{i-1}} + \frac{\partial \mathbf{x}^L}{\partial b}\frac{\partial b}{\partial \mathbf{x}^{i-1}}, \tag{4}
$$

where the second term with $\frac{\partial \mathbf{x}^L}{\partial b}$ vanishes naturally. $\square$

IV. EXAMPLES

A. Bessel function

An Bessel function of the first kind of order ν can be computed using Taylor expansion

$$
J_\nu(z) = \sum_{n=0}^{\infty} \frac{(z/2)^\nu}{\Gamma(k+1)\Gamma(k+\nu+1)}(-z^2/4)^n \tag{5}
$$

where $\Gamma(n) = (n-1)!$ is the Gamma function. One can compute the accumulated item iteratively as $s_n = -\frac{z^2}{4}s_{n-1}$. The irreversible implementation is

```
function besselj(ν, z; atol=1e-8)
    k = 0
    s = (z/2)^ν / factorial(ν)
    out = s
    while abs(s) > atol
        k += 1
        s *= (-1) / k / (k+ν) * (z/2)^2
        out += s
    end
    out
end
```

This computational process could be diagrammatically represented as a DAG as show in Fig. 6(a). In this diagram, the data is represented as an edge, it connects at most two nodes, one generates this data and one consumes it as an input parameter. Computational graph is more likely a mathematical expression, it can not describe inplace functions or control flows conveniently because it does not have the concept of memory.

In the following, we introduce the reversible implementation and the memory oriented computational graph. Intuitively, one needs an increasing size of tape to cache the intermediate state, since one can not release state s_{n-1} directly after computing s_n [17]. To reduce the memory allocation without increasing the time complexity of program, we introduce the following reversible approximate multiplier

(a) Where a vertex (circle) is an operation and a directed edge is a variable. The gray region is the body of while loop. The while loop is unrolled.

(b) Vertical lines are variables, boxes and dots connected by parallel lines are operations. A dot represents using a variable as control while a box represents mutating a variable. The gray box represents the body of a while loop, where a precondition and postcondition are positioned on top and bottom of this box.

Figure 6: The comparison between (a) tranditional and (b) memory oriented computational graph of the first kind Bessel function.

```
1  @i @inline function imul(out!, x, anc!)
2      anc! += out! * x
3      out! -= anc! / x
4      SWAP(out!, anc!)
5  end
```

Here, the definition of SWAP instruction can be found in Appendix B, anc! ≈ 0 is a *dirty ancilla*. Line 2 computes the result and accumulates it to the dirty ancilla, we get an approximately correct output in anc!. Line 3 "uncomputes" out! approximately by using the information stored in anc!, leaving a dirty zero state in register out!. Line 4 swaps the contents in out! and anc!. Finally, we have an approximately correct output and a dirtier ancilla. With this multiplier, we implementation J_ν as follows.

```
using NiLang, NiLang.AD

@i function ibesselj(out!, ν, z; atol=1e-8)
    k ← 0
    fact_nu ← zero(ν)
    halfz ← zero(z)
    halfz_power_nu ← zero(z)
    halfz_power_2 ← zero(z)
    out_anc ← zero(z)
    anc1 ← zero(z)
    anc2 ← zero(z)
    anc3 ← zero(z)
    anc4 ← zero(z)
    anc5 ← zero(z)

    @routine begin
        halfz += z / 2
        halfz_power_nu += halfz ^ ν
        halfz_power_2 += halfz ^ 2
        ifactorial(fact_nu, ν)

        anc1 += halfz_power_nu/fact_nu
        out_anc += identity(anc1)
        while (abs(unwrap(anc1)) > atol &&
                abs(unwrap(anc4)) < atol, k!=0)
            k += identity(1)
            @routine begin
                anc5 += identity(k)
                anc5 += identity(ν)
                anc2 -= k * anc5
                anc3 += halfz_power_2 / anc2
            end
            imul(anc1, anc3, anc4)
            out_anc += identity(anc1)
            ~@routine
        end
    end
    out! += identity(out_anc)
```

Here, only a constant number of ancillas are used in this implementation, while the algorithm complexity does not increase comparing to its irreversible counterpart. ancilla `anc4` plays the role of *dirty ancilla* in multiplication, it is uncomputed rigorously in the uncomputing stage. The reason why the "approximate uncomputing" trick works here lies in the fact that from the mathematic perspective the state in nth step $\{s_n, z\}$ contains the same amount of information as the state in the $n-1$th step $\{s_{n-1}, z\}$ except some special points, it is highly possible to find an equation to uncompute the previous state from the current state. This trick can be used extensively in many other application. It mitigated the artificial irreversibility brought by the number system that we have adopt at the cost of precision.

This reversible program can be diagramatically represented as a memory oriented computational graph as show in Fig. 6 (b). The key difference is we adopt a variable centric view, while a variable can be used by multiple operations. If a function uses a variable but does not change the contents of it, we call this variable a control parameter in this function. It is represented as a dot in the graph. Otherwise if it is changed, it is represented as a square. This diagram can be used to analyse uncomputable variables. In this example routine "B" changes `anc2`, `anc3` and `anc5`. while the following operation `imul` does not change these variables. Hence we can apply the inverse routine `~B` to safely restore contents `anc2`, `anc3` and `anc5`, this is what we called compute-copy-uncompute paradigm.

To obtain gradients, one can wrap the variable `out!` with `Loss` and feed it into `ibesselj'`

```julia
julia> out!, x = 0.0, 3.0
(0.0, 3.0)

julia> ibesselj'(Loss(out!), 2, x)
(Loss(GVar(0.0, 1.0)), 2, GVar(3.0, 0.0149981304))
```

Here, `ibesselj'` is a callable instance of type `Grad{typeof(ibesselj)}`. Its implementation is shown in Sec. III A. This function itself is reversible and differentiable, one can back-propagate this function to obtain Hessians as introduced in Sec. III B 1. In NiLang, it is implemented as `hessian_repeat`.

```julia
julia> hessian_repeat(ibesselj, (Loss(out!), 2, x))
3×3 Array{Float64,2}:
 0.0  0.0   0.0
 0.0  0.0   0.0
 0.0  0.0  -0.27505
```

To obtain Hessians, we can also use the Hessian propagation approach as introduced in Sec. III B 2.

```julia
julia> ibesselj''(Loss(out!), 2, x)
(Loss(BeijingRing{Float64}(0.0, 1.0, 1)), 2,
BeijingRing{Float64}(3.0, 0.014998134978750133, 2))

julia> collect_hessian()
2×2 Array{Float64,2}:
 0.0   0.0
 0.0  -0.27505
```

`ibesselj''` computes the second-order gradients. It wraps variables with type `BeijingRing` [41] in the backward pass. `BeijingRing` records Jacobians and Hessians for a variable, where Hessians are stored in a global storage. Whenever an n-th variable or ancilla is created, we push a ring of size $2n-1$ to a global tape. Whenever an ancilla is deallocated, we pop a ring from the top. The n-th ring stores Hessian elements $H_{i \leq n, n}$ and $H_{n, i < n}$. The final result can be collected by calling `collect_hessian()`, which will read out the Hessian matrix that stored in the global storage.

1. Benchmark

We have different source to source automatic different ion implementations of the first type Bessel function bench-marked and show the results below.

	Tangent/Adjoint	$T_{\min}$/**ns**	Space/KB
Julia	-	22	0
NiLang	-	59	0
ForwardDiff	Tangent	35	0
Manual	Adjoint	83	0
NiLang.AD	Adjoint	213	0
NiLang.AD (GPU)	Adjoint	1.4	0
Zygote	Adjoint	31201	13.47
Tapenade	Adjoint	?	?

Table II: Time and space used in computing the gradients of Bessel functions $J_2(1.0)$.

Julia is the CPU time used for running the irreversible forward program, is the baseline of benchmarking. NiLang is the reversible implementation, it is 2.7 times slower than its irreversible counterpart. Here, we have remove the reversibility check. ForwardDiff gives the best performance because it is designed for functions with single input. It is even faster than manually derived gradients

$$J'_\nu(z) = \frac{J_{\nu-1} - J_{\nu+1}}{2} \tag{6}$$

NiLang.AD is the reversible differential programming implementation, it considers only the backward pass. The

benchmark of its GPU version is estimated on Nvidia Titan V by broadcasting the gradient function on CUDA array of size 2^17 and take average. The Zygote benchmark considers both forward pass and backward pass. Tapenade [**JG: needs your effort, @MLS**]

```julia
@i function dot(r::T, A::SparseMatrixCSC{T},
        B::SparseMatrixCSC{T}) where {T}
    m ← size(A, 1)
    n ← size(A, 2)
    @invcheckoff branch_keeper ← zeros(Bool, 2*m)
    @safe size(B) == (m,n) ||
        throw(DimensionMismatch(
        "matrices must have the same dimensions"))
    @invcheckoff @inbounds for j = 1:n
        ia1 ← A.colptr[j]
        ib1 ← B.colptr[j]
        ia2 ← A.colptr[j+1]
        ib2 ← B.colptr[j+1]
        ia ← ia1
        ib ← ib1
        @inbounds for i=1:ia2-ia1+ib2-ib1-1
            ra ← A.rowval[ia]
            rb ← B.rowval[ib]
            if (ra == rb, ~)
                r += A.nzval[ia]' * B.nzval[ib]
            end
            # b move -> true, a move -> false
            branch_keeper[i] ⊻= ia == ia2-1 ||
                ra > rb
            ra → A.rowval[ia]
            rb → B.rowval[ib]
            if (branch_keeper[i], ~)
                ib += identity(1)
            else
                ia += identity(1)
            end
        end

        ~@inbounds for i=1:ia2-ia1+ib2-ib1-1
            # b move -> true, a move -> false
            branch_keeper[i] ⊻= ia == ia2-1 ||
                    A.rowval[ia] > B.rowval[ib]
            if (branch_keeper[i], ~)
                ib += identity(1)
            else
                ia += identity(1)
            end
        end
    end
    @invcheckoff branch_keeper → zeros(Bool, 2*m)
end
```

B. Sparse Matrices

Source to source automatic differentiation is very useful in differentiating sparse matrices. Sparse matrices can not easily benefit generic backward rules that defined for dense matrix, because generic rules does not keep the sparse structure.

Here, we show a simple example that uses reversible programming to differentiate the Frobenius dot/inner product between two sparse matrices that defined as `trace(A'B)`. This code is adapted from the irreversible implementation in Julia package `SparseArrays`.

With simple adaptation, the code becomes reversible. Here, the key point is using a `branch_keeper` vector to cache branch decisions. From the benchmark, we show time costs is approximately 1.6 times the Julia native forward implementation.

```
julia> @benchmark SparseArrays.dot($a, $b)
BenchmarkTools.Trial:
  memory estimate:  0 bytes
  allocs estimate:  0
  --------------
  minimum time:     94.537 μs (0.00% GC)
  median time:      96.959 μs (0.00% GC)
  mean time:        98.188 μs (0.00% GC)
  maximum time:     189.291 μs (0.00% GC)
  --------------
  samples:          10000
  evals/sample:     1

julia> out! = SparseArrays.dot(a, b)
25.19659286755114

julia> @benchmark (~dot)($(GVar(out!, 1.0)),
          $(GVar.(a)), $(GVar.(b)))
BenchmarkTools.Trial:   #$ TODO: remove this comment
  memory estimate:  2.17 KiB
  allocs estimate:  2
  --------------
  minimum time:     151.392 μs (0.00% GC)
  median time:      153.153 μs (0.00% GC)
  mean time:        155.884 μs (0.00% GC)
  maximum time:     270.534 μs (0.00% GC)
  --------------
  samples:          10000
  evals/sample:     1
```

Here, we turn off the irreversibility check off to achieve best performance. One can always turn off this check after the debugging. It is also possible to implement sparse matrix multiplication and other operations reversibly.

C. Unitary Matrices

A unitary matrices features uniform eigenvalues and reversibility. It is widely used as an approach to ease the gradient exploding and vanishing problem [42–44] and the memory wall problem [9]. One of the simplest way to parametrize a unitary matrix is representing a unitary matrix as a product of two-level unitary operations [44]. A real unitary matrix of size N can be parametrized compactly by $N(N-1)/2$ rotation operations [45]

$$\text{ROT}(a!, b!, \theta) = \begin{pmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{pmatrix} \begin{pmatrix} a! \\ b! \end{pmatrix}, \quad (7)$$

where θ is the rotation angle, a! and b! are target registers.

```
using NiLang, NiLang.AD

@i function umm!(x!, θ)
    @safe @assert length(θ) ==
            length(x!)*(length(x!)-1)/2
    k ← 0
    for j=1:length(x!)
        for i=length(x!)-1:-1:j
            k += identity(1)
            ROT(x![i], x![i+1], θ[k])
        end
    end

    k → length(θ)
end
```

Here, the ancilla k is deallocated manually by specifying its value, because we know the loop size is $N(N-1)/2$. We define the test functions in order to check gradients.

```
julia> @i function isum(out!, x::AbstractArray)
           for i=1:length(x)
               out! += identity(x[i])
           end
       end

julia> @i function test!(out!, x!::Vector, θ::Vector)
           umm!(x!, θ)
           isum(out!, x!)
       end

julia> out, x, θ = Loss(0.0), randn(4), randn(6);

julia> @instr test!'(out, x, θ)

julia> x
4-element Array{GVar{Float64,Float64},1}:
 GVar(1.220182125326287, 0.14540743042341095)
 GVar(2.1288634811475937, -1.3749962375499805)
 GVar(1.2696579252569677, 1.42868739498625)
 GVar(0.1083891125379283, 0.2170123344615735)

julia> @instr (~test!')(out, x, θ)

julia> x
4-element Array{Float64,1}:
 1.220182125326287
 2.1288634811475933
 1.2696579252569677
 0.10838911253792821
```

In the above testing code, test' attaches a gradient field to each element of x. ~test' is the inverse program that erase the gradient fields. Notably, this reversible implementation costs zero memory allocation although it changes the target variables inplace.

D. QR decomposition

Let's consider a naive implementation of QR decomposition from scratch. We admit this implementation is just a proof of principle which does not even consider reorthogonalization.

Figure 7: The Petersen graph with 10 vertices and 15 edges.

```
using NiLang, NiLang.AD

@i function qr(Q, R, A::Matrix{T}) where T
    anc_norm ← zero(T)
    anc_dot ← zeros(T, size(A,2))
    ri ← zeros(T, size(A,1))
    for col = 1:size(A, 1)
        ri .+= identity.(A[:,col])
        for precol = 1:col-1
            dot(anc_dot[precol], Q[:,precol], ri)
            R[precol,col] +=
                identity(anc_dot[precol])
            for row = 1:size(Q,1)
                ri[row] -=
                    anc_dot[precol] * Q[row, precol]
            end
        end
        norm2(anc_norm, ri)

        R[col, col] += anc_norm^0.5
        for row = 1:size(Q,1)
            Q[row,col] += ri[row] / R[col, col]
        end

        ~begin
            ri .+= identity.(A[:,col])
            for precol = 1:col-1
                dot(anc_dot[precol], Q[:,precol], ri)
                for row = 1:size(Q,1)
                    ri[row] -= anc_dot[precol] *
                        Q[row, precol]
                end
            end
            norm2(anc_norm, ri)
        end
    end
end
```

```
julia> A, q, r = randn(4,4), zero(A), zero(A);

julia> @i function test1(out, q, r, A)
           qr(q, r, A)
           isum(out, q)
       end

julia> check_grad(test1, (Loss(0.0), q, r, A))
true
```

Here, the loss function `test1` is defined as the sum of the output unitary matrix `q`. The `check_grad` function is a gradient checker function defined in module NiLang.AD.

E. Solving a graph embedding problem

Graph embedding has an interesting application of finding a representation for an order parameter [46] in condensed matter physics. Ref. 46 considers a problem of finding the minimum Euclidean space dimension k that a Petersen graph can fit into, with extra requirements that the distance between a pair of connected vertices has the same value l_1, the distance between a pair of disconnected vertices has the same value l_2 and $l_2 > l_1$. The Petersen graph is 10 vertices graph as shown in Fig. 7. Let's denote the set of connected and disconnected vertex pairs as L_1 and L_2 respectively. This problem can be variationally solved by differential programming by designing the following loss.

$$
\begin{aligned}
\mathcal{L} &= \mathrm{Var}(\mathrm{dist}(L_1)) + \mathrm{Var}(\mathrm{dist}(L_2)) \\
&+ \exp(\mathrm{relu}(\mathrm{E}(\mathrm{dist}(L_1)) - \mathrm{E}(\mathrm{dist}(L_2)) + 0.1))) - 1
\end{aligned} \tag{8}
$$

The first line is a summation of distance variances in two sets of vertex pairs. The second line is used to guarantee $l_2 > l_1$.

Here, in order to avoid frequent uncomputing, we allocate ancillas `ri` and `anc_dot` as vectors. The expression in ~ is used to uncompute `ri`, `anc_dot` and `anc_norm`. `dot` and `norm2` are reversible functions to compute dot product and vector norm. Their implementations could be found in Appendix D. One can easily check the correctness of the gradient function

We repeat the training for each dimension k from 1 to 9 and search for possible solutions by variationally optimizing the positions of vertices. In each training, we fixed two of the vertices and train the rest, otherwise the program will find the trivial solution with overlapped vertices. The training employs the Adam optimizer with a learning rate 0.01 [47], a 10000 steps training costs less than 1 second on a laptop. For $k = 5$, we can get a loss close to machine precision with high probability, while for $k < 5$, the loss is always much higher than 0. From the solution, it is easy to verify $l_2/l_1 = \sqrt{2}$.

V. DISCUSSION AND OUTLOOK

In this paper, we show a program on an reversible Turing machine can be differentiated to any order reliably and efficiently without sophisticated designs to memorize computational graph and intermediate states. We introduce a reversible Julia eDSL NiLang that implements a reversible AD. In a reversible programming language, we proposed to use "approximate uncomputing" trick to avoid the overhead of a reversible program in many practical cases.

In the following, we discussed some practical issues about reversible programming, and several future directions to go.

A. Time Space Tradeoff

In history, there has been many other interesting designs of reversible languages. However, current popular programming languages are all irreversible. In the simplest g-segment trade off scheme [48, 49], a RTM model has either a space overhead that proportional to computing time T or a computational overhead that sometimes can be exponential comparing with a irreversible counter part. The tradeoff between space and time is a crucial issue in the theory of RTM. In the following, we try to convince the readers that the overhead of reversible computing is not as terrible as people thought.

The overhead of reversing a program is bounded the checkpointing [50] strategy used in a traditional machine learning package that memorizes every inputs of primitives because similar strategy can also be used in reversible programming. [17] Reversible programming simply provides more alternatives to reduce the overhead. For example, the overhead in many iterative algorithms can often be removed with "arithmetic uncomputing" trick without sacrificing reversibility as shown in the `ibesselj` example in Sec. IV A.

Clever compiling can also be used to remove most overheads. Often, when we define a new reversible function, we allocate some ancillas at the beginning of the function and deallocate them through uncomputing at the end. The overhead comes from the uncomputing, in the worst case, the time used for uncomputing can be the same as the forward pass. In a hierarchical design, uncomputing can appear in every layer of abstraction. To quantify the overhead of uncomputing, we introducing the concept

***Definition* 1** (program granularity). The log ratio between

```julia
using NiLang, NiLang.AD

@i function sqdistance(dist!, x1::AbstractVector{T},
        x2::AbstractVector) where T
    @inbounds for i=1:length(x1)
        x1[i] -= identity(x2[i])
        dist! += x1[i] ^ 2
        x1[i] += identity(x2[i])
    end
end

# bonds of a petersen graph
const L1 = [(1, 6), (2, 7), (3, 8), (4, 9), (5, 10),
    (1, 2), (2, 3), (3, 4), (4, 5), (1, 5), (6, 8),
    (8, 10), (7, 10), (7, 9), (6, 9)]

# disconnected bonds of a petersen graph
const L2 = [(1, 3), (1, 4), (1, 7), (1, 8), (1, 9),
    (1, 10), (2, 4), (2, 5), (2, 6), (2, 8), (2, 9),
    (2, 10), (3, 5), (3, 6), (3, 7), (3, 9), (3, 10),
    (4, 6), (4, 7), (4, 8), (4, 10), (5, 6), (5, 7),
    (5, 8), (5, 9), (6, 7), (6, 10), (7, 8), (8, 9),
    (9, 10)]

@i function embedding_loss(out!::T, x) where T
    v1 ← zero(T)
    m1 ← zero(T)
    v2 ← zero(T)
    m2 ← zero(T)
    diff ← zero(T)
    d1 ← zeros(T, length(L1))
    d2 ← zeros(T, length(L2))
    @routine begin
        for i=1:length(L1)
            sqdistance(d1[i],
                x[:,L1[i][1]],x[:,L1[i][2]])
        end
        for i=1:length(L2)
            sqdistance(d2[i],
                x[:,L2[i][1]],x[:,L2[i][2]])
        end
        var_and_mean_sq(v1, m1, d1)
        var_and_mean_sq(v2, m2, d2)
        m1 -= identity(m2)
        m1 += identity(0.1)
    end
    out! += identity(v1)
    out! += identity(v2)
    if (m1 > 0, ~)
        # to ensure mean(v2) > mean(v1)
        # if mean(v1)+0.1 - mean(v2) > 0, punish it.
        out! += exp(m1)
        out! -= identity(1)
    end
    ~@routine
end
```

One can access the gradient by typing

```julia
iloss'(Loss(0.0), randn(5, 10))
```

the execution time of a reversible program and its irreversible counter part

The computing time increases exponentially as the granularity increases. A cleverer compilation of a program can reduce the granularity by merging the uncomputing statements to avoid repeated efforts.

At last, making reversible programming an eDSL rather than an independent language allows flexible choices between reversibility and computational overhead. For example, in order to deallocate the gradient memory in a reversible language one has to uncompute the whole process of obtaining this gradient. In our eDSL, we can just deallocate the memory irreversibly, i.e. trade energy with time.

B. Instructions and Hardwares

So far, our eDSL is compiled to Julia. In the future, it can be compiled to reversible instructions [51] and executed on a reversible device. However, arithmetic instructions should be redesigned to support better reversible programs. The major obstacle to exact reversibility programming is current floating point adders and multipliers used in our computing devices are not exactly reversible. There are proposals of reversible floating point adders and multipliers [52–55], however these designs with allocate garbage bits in each operation. Alternatives include fixed point numbers [56] and logarithmic numbers [57, 58], where logarithmic number system is reversible under *= and /=. With these infrastructures, a reversible program can be executed without suffering from the rounding error.

Reversible instructions can be executed on an energy efficient reversible hardware. In the introduction, we mentioned several reversible hardwares. A reversible hardware can be those supporting reversible gates such as the Toffoli gate and the Fredkin gate, or like an adiabatic CMOS with the ability to recover signal energy. The latter is known as the generalized reversible computing. [59, 60] In the near future, there might be energy efficient artificial intelligence (AI) chips as coprocessors that our eDSL can compile to. Since reversible computing is mainly driven by quantum computing in recent years. In the following, we comment briefly on quantum devices.

1. Quantum Computers

Building a universal quantum computer [61] is difficult. The difficulty lies in the fact that it is extremely hard to protect a quantum state. Unlike a classical state, an quantum state is can not be cloned, meanwhile, it losses information by interacting with the environment, or decoherence. Classical reversible computing does not enjoy the quantum advantage, nor the quantum disadvantages of non-cloning and decoherence. It is technically more smooth to have a reversible computing device to bridge the gap between classical devices and universal quantum computing devices. By introducing entanglement little by little, we can accelerate some basic components in reversible computing. For example, quantum Fourier transformation provides an interesting alternative to the reversible adders and multipliers by introducing the CPHASE quantum gate. [62] The development of reversible compiling theory can be benefit quantum compiling directly.

C. Outlook

The reversible eDSL NiLang can be used to solve many existing scientific computing problems. First of all, it can be used to generate AD rules for existing machine learning packages like Zygote. For example, one can use NiLang to generate backward rules for singular value decomposition and eigenvalue decomposition functions that extensively used in scientific computing [63, 64]. Although their backward rules [65–67] have been derived in recent years, these backward rules can not handle degenerate eigenvalues properly. Hopefully, the automatically generated backward rules do not have such problems.

Secondly, we can use it to overcome the memory wall problem in some applications. NiLang provides a systematic time-space trade off scheme through uncomputing. A successful related example is the memory efficient domain-specific AD engine in quantum simulator Yao [9]. This domain-specific AD engine is written in a reversible style and solved the memory bottleneck in variational quantum simulations. It also gives hitherto the best performance in differentiating quantum circuit parameters. Similarity, we can write memory efficient normalizing flow [68] with NiLang. Normalizing flow is a successful class of generative model in both computer vision [69] and quantum physics [70, 71], where its building block bijector is reversible. We can use similar idea to differentiate reversible integrators [72, 73]. With reversible integrators, it should be possible to rewrite the control system in robotics [74] in a reversible style, where scalars are first class citizen rather than tensors. Writing a control program reversibly should boost the training performance a lot.

Thirdly, reversibility is a resource for training. For those who are interested in non-gradient based training. In Appendix C, we provide a self-consistency training strategy for reversible programs.

Lastly, the reversible IR is a good starting point to study quantum compiling. Most quantum programming language preassumes a classical coprocessor and use classical control flows [75] in universal quantum computing. However, we believe reversible control flows are also very important to a universal quantum computer.

To solve the above problems better, NiLang can be improved from multiple perspectives. Like we need a more efficient fixed point or log number system to avoid rounding errors. Currently we simulate reversible arithmetics with the julia fixed point number package. [56] Then we should optimize the compiling to decreases granularity of a program and reduces uncomputing overheads. There are also some known issues to be solved like the type inference problem, we have listed some of them on GitHub. These improvements

needs participation of people from multiple fields.

[1] L. Hascoet and V. Pascual, ACM Transactions on Mathematical Software (TOMS) **39**, 20 (2013).

[2] A. Paszke, S. Gross, S. Chintala, G. Chanan, E. Yang, Z. De-Vito, Z. Lin, A. Desmaison, L. Antiga, and A. Lerer, in *NIPS Autodiff Workshop* (2017).

[3] M. Innes, E. Saba, K. Fischer, D. Gandhi, M. C. Rudilosso, N. M. Joy, T. Karmali, A. Pal, and V. Shah, "Fashionable modelling with flux," (2018), arXiv:1811.01457 [cs.PL].

[4] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefow-icz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke, Y. Yu, and X. Zheng, "TensorFlow: Large-scale machine learning on heterogeneous systems," (2015), software available from tensorflow.org.

[5] M. Innes, "Don't unroll adjoint: Differentiating ssa-form programs," (2018), arXiv:1810.07951 [cs.PL].

[6] M. Innes, A. Edelman, K. Fischer, C. Rackauckas, E. Saba, V. B. Shah, and W. Tebbutt, CoRR **abs/1907.07587** (2019), arXiv:1907.07587.

[7] Z. C. Lipton, J. Berkowitz, and C. Elkan, "A critical review of recurrent neural networks for sequence learning," (2015), arXiv:1506.00019 [cs.LG].

[8] K. He, X. Zhang, S. Ren, and J. Sun, 2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR) (2016), 10.1109/cvpr.2016.90.

[9] X.-Z. Luo, J.-G. Liu, P. Zhang, and L. Wang, "Yao.jl: Extensible, efficient framework for quantum algorithm design," (2019), arXiv:1912.10877 [quant-ph].

[10] M. MacKay, P. Vicol, J. Ba, and R. B. Grosse, in *Advances in Neural Information Processing Systems 31*, edited by S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett (Curran Associates, Inc., 2018) pp. 9029–9040.

[11] J. Behrmann, D. Duvenaud, and J. Jacobsen, CoRR **abs/1811.00995** (2018), arXiv:1811.00995.

[12] D. Maclaurin, D. Duvenaud, and R. Adams, in *Proceedings of the 32nd International Conference on Machine Learning*, Proceedings of Machine Learning Research, Vol. 37, edited by F. Bach and D. Blei (PMLR, Lille, France, 2015) pp. 2113–2122.

[13] A. N. Gomez, M. Ren, R. Urtasun, and R. B. Grosse, in *Advances in Neural Information Processing Systems 30*, edited by I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett (Curran Associates, Inc., 2017) pp. 2214–2224.

[14] J.-H. Jacobsen, A. W. Smeulders, and E. Oyallon, in *International Conference on Learning Representations* (2018).

[15] J. Bezanson, S. Karpinski, V. B. Shah, and A. Edelman, arXiv preprint arXiv:1209.5145 (2012).

[16] J. Bezanson, A. Edelman, S. Karpinski, and V. B. Shah, SIAM Review **59**, 65–98 (2017).

[17] K. S. Perumalla, *Introduction to reversible computing* (Chapman and Hall/CRC, 2013).

[18] M. P. Frank, IEEE Spectrum **54**, 32–37 (2017).

[19] C. Lutz, "Janus: a time-reversible language," (1986), *Letter to R. Landauer*.

[20] M. P. Frank, *The R programming language and compiler*, Tech. Rep. (MIT Reversible Computing Project Memo, 1997).

[21] I. Lanese, N. Nishida, A. Palacios, and G. Vidal, Journal of Logical and Algebraic Methods in Programming **100**, 71–97 (2018).

[22] T. Haulund, "Design and implementation of a reversible object-oriented programming language," (2017), arXiv:1707.07845 [cs.PL].

[23] M. P. Frank and T. F. Knight Jr, *Reversibility for efficient computing*, Ph.D. thesis, Massachusetts Institute of Technology, Dept. of Electrical Engineering and . . . (1999).

[24] J. G. Koller and W. C. Athas, in *Workshop on Physics and Computation* (1992) pp. 267–270.

[25] R. C. Merkle, R. A. Freitas, T. Hogg, T. E. Moore, M. S. Moses, and J. Ryley, Journal of Mechanisms and Robotics **10** (2018), 10.1115/1.4041209.

[26] K. Likharev, IEEE Transactions on Magnetics **13**, 242 (1977).

[27] V. K. Semenov, G. V. Danilov, and D. V. Averin, IEEE Transactions on Applied Superconductivity **13**, 938 (2003).

[28] R. Landauer, IBM journal of research and development **5**, 183 (1961).

[29] E. P. DeBenedictis, J. K. Mee, and M. P. Frank, Computer **50**, 76 (2017).

[30] D. R. Jefferson, ACM Transactions on Programming Languages and Systems (TOPLAS) **7**, 404 (1985).

[31] B. Boothe, ACM SIGPLAN Notices **35**, 299 (2000).

[32] G. H. Golub and C. F. Van Loan, *Matrix computations*, Vol. 3 (JHU press, 2012).

[33] C. H. Bennett (1973).

[34] "MLStyle.jl," https://github.com/thautwarm/MLStyle.jl.

[35] Wikipedia contributors, "Einstein notation — Wikipedia, the free encyclopedia," (2020), [Online; accessed 20-February-2020].

[36] R. Orús, Annals of Physics **349**, 117–158 (2014).

[37] J. Revels, M. Lubin, and T. Papamarkou, "Forward-mode automatic differentiation in julia," (2016), arXiv:1607.07892 [cs.MS].

[38] A. McInerney, *First steps in differential geometry* (Springer, 2015).

[39] "Paper's Github Repository," https://github.com/GiggleLiu/nilangpaper/tree/master/codes.

[40] J. Martens, I. Sutskever, and K. Swersky, "Estimating the hessian by back-propagating curvature," (2012), arXiv:1206.6464 [cs.LG].

[41] When people ask for the location in Beijing, they will start by asking which ring it is? We use the similar approach to locate the elements of Hessian matrix.

[42] M. Arjovsky, A. Shah, and Y. Bengio, CoRR **abs/1511.06464** (2015), arXiv:1511.06464.

[43] S. Wisdom, T. Powers, J. R. Hershey, J. L. Roux, and L. Atlas, "Full-capacity unitary recurrent neural networks," (2016), arXiv:1611.00035 [stat.ML].

[44] L. Jing, Y. Shen, T. Dubcek, J. Peurifoy, S. A. Skirlo, M. Tegmark, and M. Soljacic, CoRR **abs/1612.05231** (2016), arXiv:1612.05231.

[45] C.-K. LI, R. ROBERTS, and X. YIN, International Journal of Quantum Information **11**, 1350015 (2013).

[46] J. Takahashi and A. W. Sandvik, "Valence-bond solids, vestigial order, and emergent so(5) symmetry in a two-dimensional quantum magnet," (2020), arXiv:2001.10045 [cond-mat.str-el].

[47] D. P. Kingma and J. Ba, arXiv:1412.6980.

[48] C. H. Bennett, SIAM Journal on Computing **18**, 766 (1989).

[49] R. Y. Levine and A. T. Sherman, SIAM Journal on Computing **19**, 673 (1990).

[50] T. Chen, B. Xu, C. Zhang, and C. Guestrin, CoRR **abs/1604.06174** (2016), arXiv:1604.06174.

[51] C. J. Vieri, *Reversible Computer Engineering and Architecture*, Ph.D. thesis, Cambridge, MA, USA (1999), aAI0800892.

[52] M. Nachtigal, H. Thapliyal, and N. Ranganathan, in *10th IEEE International Conference on Nanotechnology* (2010) pp. 233–237.

[53] M. Nachtigal, H. Thapliyal, and N. Ranganathan, in *2011 11th IEEE International Conference on Nanotechnology* (2011) pp. 451–456.

[54] T. D. Nguyen and R. V. Meter, "A space-efficient design for reversible floating point adder in quantum computing," (2013), arXiv:1306.3760 [quant-ph].

[55] T. Häner, M. Soeken, M. Roetteler, and K. M. Svore, "Quantum circuits for floating-point arithmetic," (2018), arXiv:1807.02023 [quant-ph].

[56] "FixedPointNumbers.jl," https://github.com/JuliaMath/FixedPointNumbers.jl.

[57] F. J. Taylor, R. Gill, J. Joseph, and J. Radke, IEEE Transactions on Computers **37**, 190 (1988).

[58] "LogarithmicNumbers.jl," https://github.com/cjdoris/LogarithmicNumbers.jl.

[59] M. P. Frank, in *35th International Symposium on Multiple-Valued Logic (ISMVL'05)* (2005) pp. 168–185.

[60] M. P. Frank, in *Reversible Computation*, edited by I. Phillips and H. Rahaman (Springer International Publishing, Cham, 2017) pp. 19–34.

[61] M. A. Nielsen and I. Chuang, "Quantum computation and quantum information," (2002).

[62] L. Ruiz-Perez and J. C. Garcia-Escartin, Quantum Information Processing **16** (2017), 10.1007/s11128-017-1603-1.

[63] J.-G. L. Hao Xie and L. Wang, arXiv:2001.04121.

[64] H.-J. Liao, J.-G. Liu, L. Wang, and T. Xiang, Physical Review X **9** (2019), 10.1103/physrevx.9.031041.

[65] M. Seeger, A. Hetzel, Z. Dai, E. Meissner, and N. D. Lawrence, "Auto-differentiating linear algebra," (2017), arXiv:1710.08717 [cs.MS].

[66] Z.-Q. Wan and S.-X. Zhang, "Automatic differentiation for complex valued svd," (2019), arXiv:1909.02659 [math.NA].

[67] C. Hubig, "Use and implementation of autodifferentiation in tensor network methods with complex scalars," (2019), arXiv:1907.13422 [cond-mat.str-el].

[68] I. Kobyzev, S. Prince, and M. A. Brubaker, "Normalizing flows: An introduction and review of current methods," (2019), arXiv:1908.09257 [stat.ML].

[69] D. P. Kingma and P. Dhariwal, in *Advances in Neural Information Processing Systems 31*, edited by S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett (Curran Associates, Inc., 2018) pp. 10215–10224.

[70] L. Dinh, J. Sohl-Dickstein, and S. Bengio, "Density estimation using real nvp," (2016), arXiv:1605.08803 [cs.LG].

[71] S.-H. Li and L. Wang, Physical Review Letters **121** (2018), 10.1103/physrevlett.121.260601.

[72] P. Hut, J. Makino, and S. McMillan, The Astrophysical Journal **443**, L93 (1995).

[73] D. N. Laikov, Theoretical Chemistry Accounts **137** (2018), 10.1007/s00214-018-2344-7.

[74] M. Giftthaler, M. Neunert, M. Stäuble, M. Frigerio, C. Semini, and J. Buchli, Advanced Robotics **31**, 1225–1237 (2017).

[75] K. Svore, M. Roetteler, A. Geller, M. Troyer, J. Azariah, C. Granade, B. Heim, V. Kliuchnikov, M. Mykhailova, and A. Paz, Proceedings of the Real World Domain Specific Languages Workshop 2018 on - RWDSL2018 (2018), 10.1145/3183895.3183901.

[76] M. Bender, P.-H. Heenen, and P.-G. Reinhard, Rev. Mod. Phys. **75**, 121 (2003).

[77] T. Besard, C. Foket, and B. D. Sutter, CoRR **abs/1712.03112** (2017), arXiv:1712.03112.

Appendix A: NiLang Grammar

To define a reversible function one can use "@i" plus a normal function definition like bellow

```
"""
docstring...
"""
@i function f(args..., kwargs...) where {...}
    <stmts>
end
```

where the definition of "<stmts>" are shown in the grammar on the next column. The following is a list of terminologies used in the definition of grammar

- *ident*, symbols

- *num*, numbers

- ϵ, empty statement

- *JuliaExpr*, native Julia expression

- [ ], zero or one repetitions.

Here, all *JuliaExpr* should be pure, otherwise the reversibility is not guaranteed. Dataview is a view of a data, it can be a bijective mapping of an object, an item of an array or a field of an object.

⟨Stmts⟩ ::= ϵ
   | ⟨Stmt⟩
   | ⟨Stmts⟩ ⟨Stmt⟩
⟨Stmt⟩ ::= ⟨BlockStmt⟩
   | ⟨IfStmt⟩
   | ⟨WhileStmt⟩
   | ⟨ForStmt⟩
   | ⟨InstrStmt⟩
   | ⟨RevStmt⟩
   | ⟨AncillaStmt⟩
   | ⟨TypecastStmt⟩
   | ⟨@routine⟩ ⟨Stmt⟩
   | ⟨@safe⟩ *JuliaExpr*
   | ⟨CallStmt⟩
⟨BlockStmt⟩ ::= *begin* ⟨Stmts⟩ *end*
⟨RevCond⟩ ::= *(JuliaExpr , JuliaExpr)*
⟨IfStmt⟩ ::= *if* ⟨RevCond⟩ ⟨Stmts⟩ [*else* ⟨Stmts⟩] *end*
⟨WhileStmt⟩ ::= *while* ⟨RevCond⟩ ⟨Stmts⟩ *end*
⟨Range⟩ ::= *JuliaExpr : JuliaExpr* [*: JuliaExpr*]
⟨ForStmt⟩ ::= *for ident =* ⟨Range⟩ ⟨Stmts⟩ *end*
⟨KwArg⟩ ::= *ident = JuliaExpr*
⟨KwArgs⟩ ::= [⟨KwArgs⟩ ,] ⟨KwArg⟩
⟨CallStmt⟩ ::= *JuliaExpr (* [⟨DataViews⟩] [*;* ⟨KwArgs⟩] *)*
⟨Constant⟩ ::= *num* | π
⟨InstrBinOp⟩ ::= += | -= | ⊻=
⟨InstrTrailer⟩ ::= [.] *(* [⟨DataViews⟩] *)*
⟨InstrStmt⟩ ::= ⟨DataView⟩ ⟨InstrBinOp⟩ *ident* [⟨InstrTrailer⟩]
⟨RevStmt⟩ ::= ~ ⟨Stmt⟩
⟨AncillaStmt⟩ ::= *ident* ← *JuliaExpr*
⟨TypecastStmt⟩ ::= *(JuliaExpr => JuliaExpr) (ident)*
⟨@routine⟩ ::= *@routine ident* ⟨Stmt⟩
⟨@safe⟩ ::= *@safe JuliaExpr*
⟨DataViews⟩ ::= ϵ
   | ⟨DataView⟩
   | ⟨DataViews⟩ , ⟨DataView⟩
   | ⟨DataViews⟩ , ⟨DataView⟩ ...
⟨DataView⟩ ::= ⟨DataView⟩ *[JuliaExpr]*
   | ⟨DataView⟩ . *ident*
   | *JuliaExpr (* ⟨DataView⟩ *)*
   | ⟨DataView⟩ ′
   | - ⟨DataView⟩
   | ⟨Constant⟩
   | *ident*

Appendix B: Instructions and Backward Rules

The translation of instructions to Julia functions
The list of instructions implemented in NiLang

instruction	translated	symbol
y += $f(args...)$	PlusEq(f)(args...)	$\oplus$
y -= $f(args...)$	MinusEq(f)(args...)	$\ominus$
y $\veebar$= $f(args...)$	XorEq(f)(args...)	$\odot$

Table III: Instructions and their compilation in NiLang.

instruction	output
SWAP(a, b)	b, a
ROT(a, b, θ)	$a\cos\theta - b\sin\theta, b\cos\theta + a\sin\theta, \theta$
IROT(a, b, θ)	$a\cos\theta + b\sin\theta, b\cos\theta - a\sin\theta, \theta$
MULINT(a, b)	$a * b, b$
DIVINT(a, b)	$a/b, b$
y += $a\hat{\ }b$	$y + a^b, a, b$
y += $\exp(x)$	$y + e^x, x$
y += $\log(x)$	$y + \log x, x$
y += $\sin(x)$	$y + \sin x, x$
y += $\cos(x)$	$y + \cos x, x$
y += $\text{abs}(x)$	$y + |x|, x$
NEG(y)	$-y$
CONJ(y)	y'

Table IV: A collection of reversible instructions, "." is the broadcasting operations in Julia. The second argument of MULINT and DIVINT should be a nonzero integer.

1. Backward rules for instructions

For function $\vec{y} = f(\vec{x})$, its Jacobian is $J_{ij} = \frac{\partial y_i}{\partial x_j}$ and its Hessian is $H^k_{ij} = \frac{\partial y_k}{x_i x_j}$. We have the following local Jacobians and Hessians on the above instructions.

1. a += b

$$J = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}$$
$$H = \mathbf{0}$$

The inverse is a -= b, its Jacobian is the inverse of the matrix above

$$J(f^{-1}) = J^{-1} = \begin{pmatrix} 1 & -1 \\ 0 & 1 \end{pmatrix}$$

In the following, we omit the Jacobians and Hessians of inverse functions.

2. a += $b * c$

$$J = \begin{pmatrix} 1 & c & b \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$
$$H^a_{bc} = H^a_{cb} = 1, else\ 0$$

3. a += b/c

$$J = \begin{pmatrix} 1 & 1/c & -b/c^2 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$
$$H^a_{cc} = 2b/c^3,$$
$$H^a_{bc} = H^a_{cb} = -1/c^2, else\ 0$$

4. a += b^c

$$J = \begin{pmatrix} 1 & cb^{c-1} & b^c \log b \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$
$$H^a_{bc} = H^a_{cb} = b^{c-1} + cb^{c-1} \log b,$$
$$H^a_{bb} = (c-1)cb^{c-2},$$
$$H^a_{cc} = b^c \log^2 b, else\ 0$$

5. a += e^b

$$J = \begin{pmatrix} 1 & e^b \\ 0 & 1 \end{pmatrix}$$
$$H^a_{bb} = e^b, else\ 0$$

6. a += $\log b$

$$J = \begin{pmatrix} 1 & 1/b \\ 0 & 1 \end{pmatrix}$$
$$H^a_{bb} = -1/b^2, else\ 0$$

7. a += $\sin b$

$$J = \begin{pmatrix} 1 & \cos b \\ 0 & 1 \end{pmatrix}$$
$$H^a_{bb} = -\sin b, else\ 0$$

8. a += $\cos b$

$$J = \begin{pmatrix} 1 & -\sin b \\ 0 & 1 \end{pmatrix}$$
$$H^a_{bb} = -\cos b, else\ 0$$

9. $a \mathrel{+}= |b|$

$$J = \begin{pmatrix} 1 & \text{sign}(b) \\ 0 & 1 \end{pmatrix}$$
$$H = \mathbf{0}$$

10. $a = -a$

$$J = \begin{pmatrix} -1 \end{pmatrix}$$
$$H = \mathbf{0}$$

11. $\text{SWAP}(a, b) = (b, a)$

$$J = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$
$$H = \mathbf{0}$$

12.

$$\text{ROT}(a, b, \theta) = \begin{pmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{pmatrix} \begin{pmatrix} a \\ b \end{pmatrix}$$

$$J = \begin{pmatrix} \cos\theta & -\sin\theta & -b\cos\theta - a\sin\theta \\ \sin\theta & \cos\theta & a\cos\theta - b\sin\theta \\ 0 & 0 & 1 \end{pmatrix}$$
$$H^a_{a\theta} = H^a_{\theta,a} = -\sin\theta,$$
$$H^a_{b\theta} = H^a_{\theta,b} = -\cos\theta,$$
$$H^a_{\theta\theta} = -a\cos\theta + b\sin\theta,$$
$$H^b_{a\theta} = H^b_{\theta a} = \cos\theta,$$
$$H^b_{b\theta} = H^b_{\theta b} = -\sin\theta,$$
$$H^b_{\theta\theta} = -b\cos\theta - a\sin\theta, else\ 0$$

Appendix C: Learn by consistency

Consider a training that with input $\mathbf{x}^*$ and output $\mathbf{y}^*$, find a set of parameters $\mathbf{p}_x$ that satisfy $\mathbf{y}^* = f(\mathbf{x}^*, \mathbf{p}_x)$. In traditional machine learning, we define a loss $\mathcal{L} = \text{dist}(\mathbf{y}^*, f(\mathbf{x}^*, \mathbf{p}_x))$ and minimize it with gradient $\frac{\partial L}{\partial \mathbf{p}_x}$. This works only when the target function is locally differentiable.

Here we provide an alternative by making use of reversibility. We construct a reversible program $\mathbf{y}, \mathbf{p}_y = f_r(\mathbf{x}, \mathbf{p}_x)$, where $\mathbf{p}_x$ and $\mathbf{p}_y$ are "parameter" spaces on the input side and output

side. The algorithm can be summarized as

Algorithm 2: Learn by consistency

Result: $\mathbf{p}_x$
Initialize $\mathbf{x}$ to $\mathbf{x}^*$, parameter space $\mathbf{p}_x$ to random.
if $\mathbf{p}_y$ *is* `null` **then**
  | $\mathbf{x}, \mathbf{p}_x = f_r^{-1}(\mathbf{y}^*)$
else
  | $\mathbf{y}, \mathbf{p}_y = f_r(\mathbf{x}, \mathbf{p}_x)$
  | **while** $\mathbf{y} \not\approx \mathbf{y}^*$ **do**
  | | $\mathbf{y} = \mathbf{y}^*$
  | | $\mathbf{x}, \mathbf{p}_x = f_r^{-1}(\mathbf{y}, \mathbf{p}_y)$.
  | | $\mathbf{x} = \mathbf{x}^*$
  | | $\mathbf{y}, \mathbf{p}_y = f_r(\mathbf{x}, \mathbf{p}_x)$

Here, `parameter`$(\cdot)$ is a function for taking the parameter space. This algorithm utilizes the self-consistency relation

$$\mathbf{p}_x^* = \texttt{parameter}(f_r^{-1}(\mathbf{y}^*, \texttt{parameter}(f_r(\mathbf{x}^*, \mathbf{p}_x^*)))), \quad \text{(C1)}$$

Similar idea of training by consistency is used in self-consistent meanfield theory [76] in physics. Finding the self-consistent relation is crucial to a self-consistency based training. Here, the reversibility provides a natural self-consistency relation. However, it is not a silver bullet, let's consider the following example

```
@i function f1(y!, x, p!)
    p! += identity(x)
    y! -= exp(x)
    y! += exp(p!)
end

@i function f2(y!, x!, p!)
    p! += identity(x!)
    y! -= exp(x!)
    x! -= log(-y!)
    y! += exp(p!)
end

function train(f)
    loss = Float64[]
    p = 1.6
    for i=1:100
        y!, x = 0.0, 0.3
        @instr f(y!, x, p)
        push!(loss, y!)
        y! = 1.0
        @instr (~f)(y!, x, p)
    end
    loss
end
```

Functions `f1` and `f2` computes $f(x, p) = e^{(p+x)} - e^x$ and stores the output in a new memory `y!`. The only difference is `f2` uncomputes x arithmetically. The task of training is to find a p that make the output value equal to target value 1. After 100 steps, `f2` runs into the fixed point with x equal

Figure 8: The output value `y!` as a function of self-consistent training step.

to 1 upto machine precision. However, parameters in `f1` does change at all. The training of `f1` fails because this function actually computes $f1(y, x, p) = y + e^{(p+x)} - e^x, x, x+p$, where the training parameter p is completely determined by the parameter space on the output side $x \cup x + p$. As a result, shifting y directly is the only approach to satisfy the consistency relation. On the other side, $f2(y, x, p) = y + e^{(p+x)} - e^x, \tilde{0}, x + p$, the output parameters $\tilde{0} \cup x + p$ can not uniquely determine input parameters p and x. Here, we use $\tilde{0}$ to denote the zero with rounding error.

By viewing **x** and parameters in $\mathbf{p}_x$ as variables, we can study the trainability from the information perspective.

Theorem 2. *Only if the the conditional entropy $S(\mathbf{y}|\mathbf{p}_y)$ is nonzero, algorithm 2 is trainable.*

Proof. The above example reveals a fact that training impossible when output parameters completely determines input parameters (or $S(\mathbf{p}_x|\mathbf{p}_y) = 0$).

$$\begin{aligned}
S(\mathbf{p}_x|\mathbf{p}_y) &= S(\mathbf{p}_x \cup \mathbf{p}_y) - S(\mathbf{p}_y) \\
&\leq S\left((\mathbf{p}_x \cup \mathbf{x}) \cup \mathbf{p}_y\right) - S(\mathbf{p}_y), \\
&\leq S\left((\mathbf{p}_y \cup \mathbf{y}) \cup \mathbf{p}_y\right) - S(\mathbf{p}_y), \\
&\leq S(\mathbf{y}|\mathbf{p}_y).
\end{aligned} \tag{C2}$$

The third line uses the bijectivity $S(\mathbf{x} \cup \mathbf{p}_x) = S(\mathbf{y} \cup \mathbf{p}_y)$. This inequality shows that when the parameter space on the output side satisfies $S(\mathbf{y}|\mathbf{p}_y) = 0$, i.e. contains all information to determine the output field, the input parameters are also completely determined by this parameter space, hence training can not work. □

In the above example, it corresponds to the case $S\left(e^{(x+y)-e^x}|x \cup x + y\right) = 0$ in $f1$. The solution is to remove the information redundancy in output parameter space through uncomputing as shown in $f2$.

Appendix D: Functions used in the main text

```
"""
    dot(out!, v1, v2)

dot product.
"""
@i function dot(out!, v1::Vector{T}, v2) where T
    for i = 1:length(v1)
        out! += v1[i]'*v2[i]
    end
end


"""
    norm2(out!, vec)

squared norm.
"""
@i function norm2(out!, vec::Vector{T}) where T
    anc1 ← zero(T)
    for i = 1:length(vec)
        anc1 += identity(vec[i]')
        out! += anc1*vec[i]
        anc1 -= identity(vec[i]')
    end
end
```

In `norm2`, we copied `vec[i]'` to `anc1` to avoid the same variable appear twice in the argument list of $\oplus(*)$, where the prime represents the adjoint dataview.

```
"""
    var_and_mean_sq(var!, mean!, sqv)

Variance and mean value from squared values `sqv`.
"""
@i function var_and_mean_sq(var!, mean!, sqv)
    sqmean ← zero(mean!)
    @inbounds for i=1:length(sqv)
        mean! += sqv[i] ^ 0.5
        var! += identity(sqv[i])
    end
    DIVINT(mean!, length(sqv))
    DIVINT(var!, length(sqv))
    sqmean += mean! ^ 2
    var! -= identity(sqmean)
    sqmean -= mean! ^ 2
    MULINT(var!, length(sqv))
    DIVINT(var!, length(sqv)-1)
end
```

Appendix E: CUDA compitibility

CUDA programming is playing a more and more important role in high performance computing. In Julia, one can write kernel functions in native Julia language with CUDA-

native [77]. NiLang is compatible with CUDAnative, one can write a reversible kernel like the following.

```julia
using CuArrays, CUDAnative, GPUArrays
using NiLang, NiLang.AD

@i @inline function swap_kernel(state, mask1, mask2)
    @invcheckoff b ← (blockIdx().x-1) *
        blockDim().x + threadIdx().x - 1
    @invcheckoff if (b < length(state), ~)
        if (b&mask1==0 && b&mask2==mask2, ~)
            SWAP(state[b+1], state[b ∨
                (mask1|mask2) + 1])
        end
    end
end
```

This kernel function simulates the SWAP gate in quantum computing. Here, we use the macro `@invcheckoff` to remove the reversibility checks. It is necessary because the possible error thrown in a kernel function can not be handled properly on a CUDA thread. One can launch this kernel function to GPUs with a single macro `@cuda` as show in the following using case.

```julia
julia> @i function instruct!(state::CuVector,
            gate::Val{:SWAP}, locs::Tuple{Int,Int})
        mask1 ← 1 << (tget(locs, 1)-1)
        mask2 ← 1 << (tget(locs, 2)-1)
        XY ← GPUArrays.thread_blocks_heuristic(
            length(state))
        @cuda threads=tget(XY,1) blocks=tget(XY,
            2) swap_kernel(state, mask1, mask2)
    end

julia> instruct!(CuArray(randn(8)),
        Val(:SWAP), (1,3))[1]
8-element CuArray{Float64,1,Nothing}:
 -0.06956048379200473
 -0.6464176838567472
 -0.06523362834285944
 -0.7314356941903547
  1.512329204247244
  0.9773772766637732
  1.6473223915215722
 -1.0631789613639087
```

Appendix F: Computing Fibonacci Numbers

The following is an example that everyone likes, computing Fibonacci number recursively.

```julia
using NiLang

@i function rfib(out!, n::T) where T
    n1 ← zero(T)
    n2 ← zero(T)
    @routine begin
        n1 += identity(n)
        n1 -= identity(1)
        n2 += identity(n)
        n2 -= identity(2)
    end
    if (value(n) <= 2, ~)
        out! += identity(1)
    else
        rfib(out!, n1)
        rfib(out!, n2)
    end
    ~@routine
end
```

The time complexity of this recursive algorithm is exponential to input `n`. It is also possible to write a reversible linear time for loop algorithm. A slightly non-trivial task is computing the first Fibonacci number that greater or equal to a certain number z, where a `while` statement is required.

```julia
@i function rfibn(n!, z)
    @safe @assert n! == 0
    out ← 0
    rfib(out, n!)
    while (out < z, n! != 0)
        ~rfib(out, n!)
        n! += identity(1)
        rfib(out, n!)
    end
    ~rfib(out, n!)
end
```

In this example, the postcondition n!=0 in the `while` statement is false before entering the loop, and becomes true in later iterations. In the reverse program, the `while` statement stops at n==0. If executed correctly, a user will see the following result.

```julia
julia> rfib(0, 10)
(55, 10)

julia> rfibn(0, 100)
(12, 100)

julia> (~rfibn)(rfibn(0, 100)...)
(0, 100)
```

This examplifies how an addition postcondition provided by user can help reversing a control flow without caching

controls.